Natural Language Processing

Lecture 11

Dirk Hovy

dirk.hovy@unibocconi.it

@dirk_hovy

Bocconi

Examples

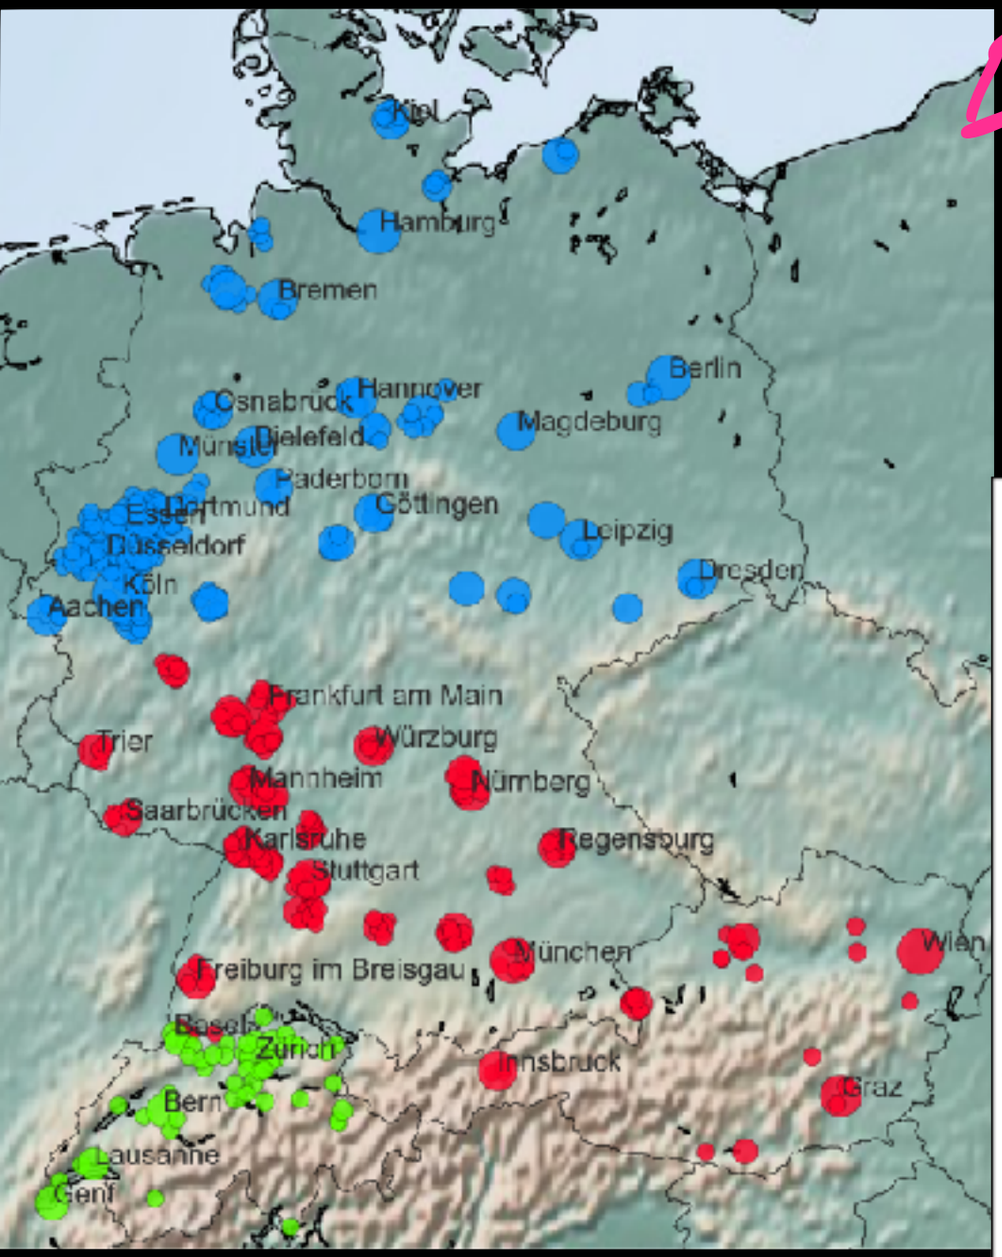

Clustering

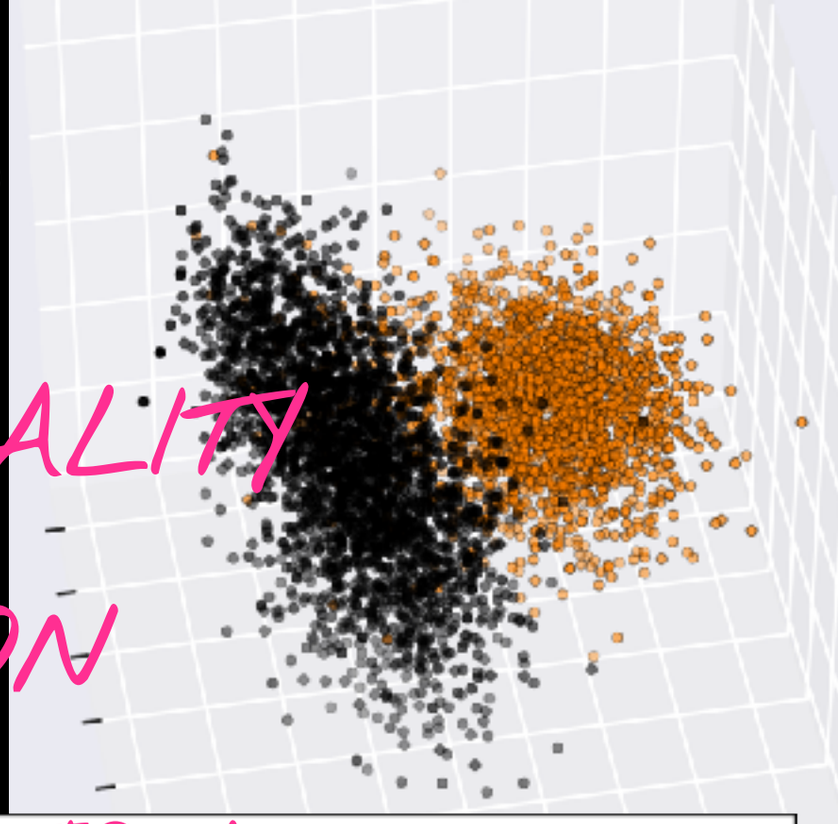

Dimensionality Reduction

Visualization

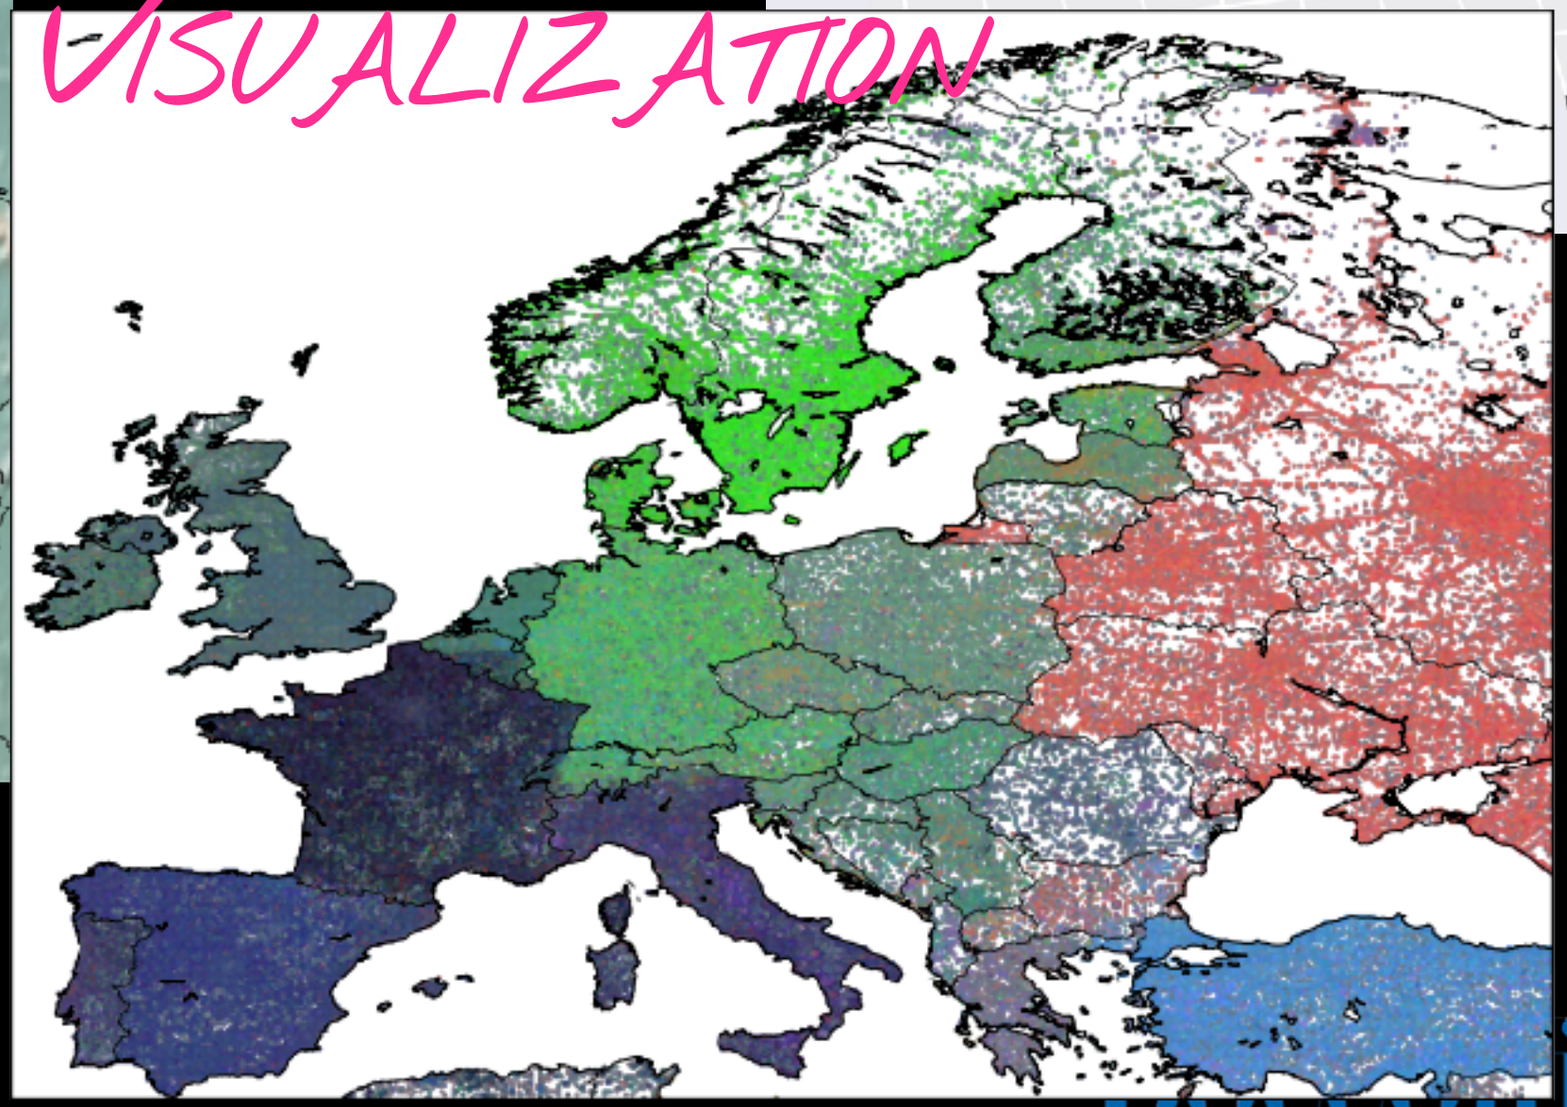

Latent Dimensions

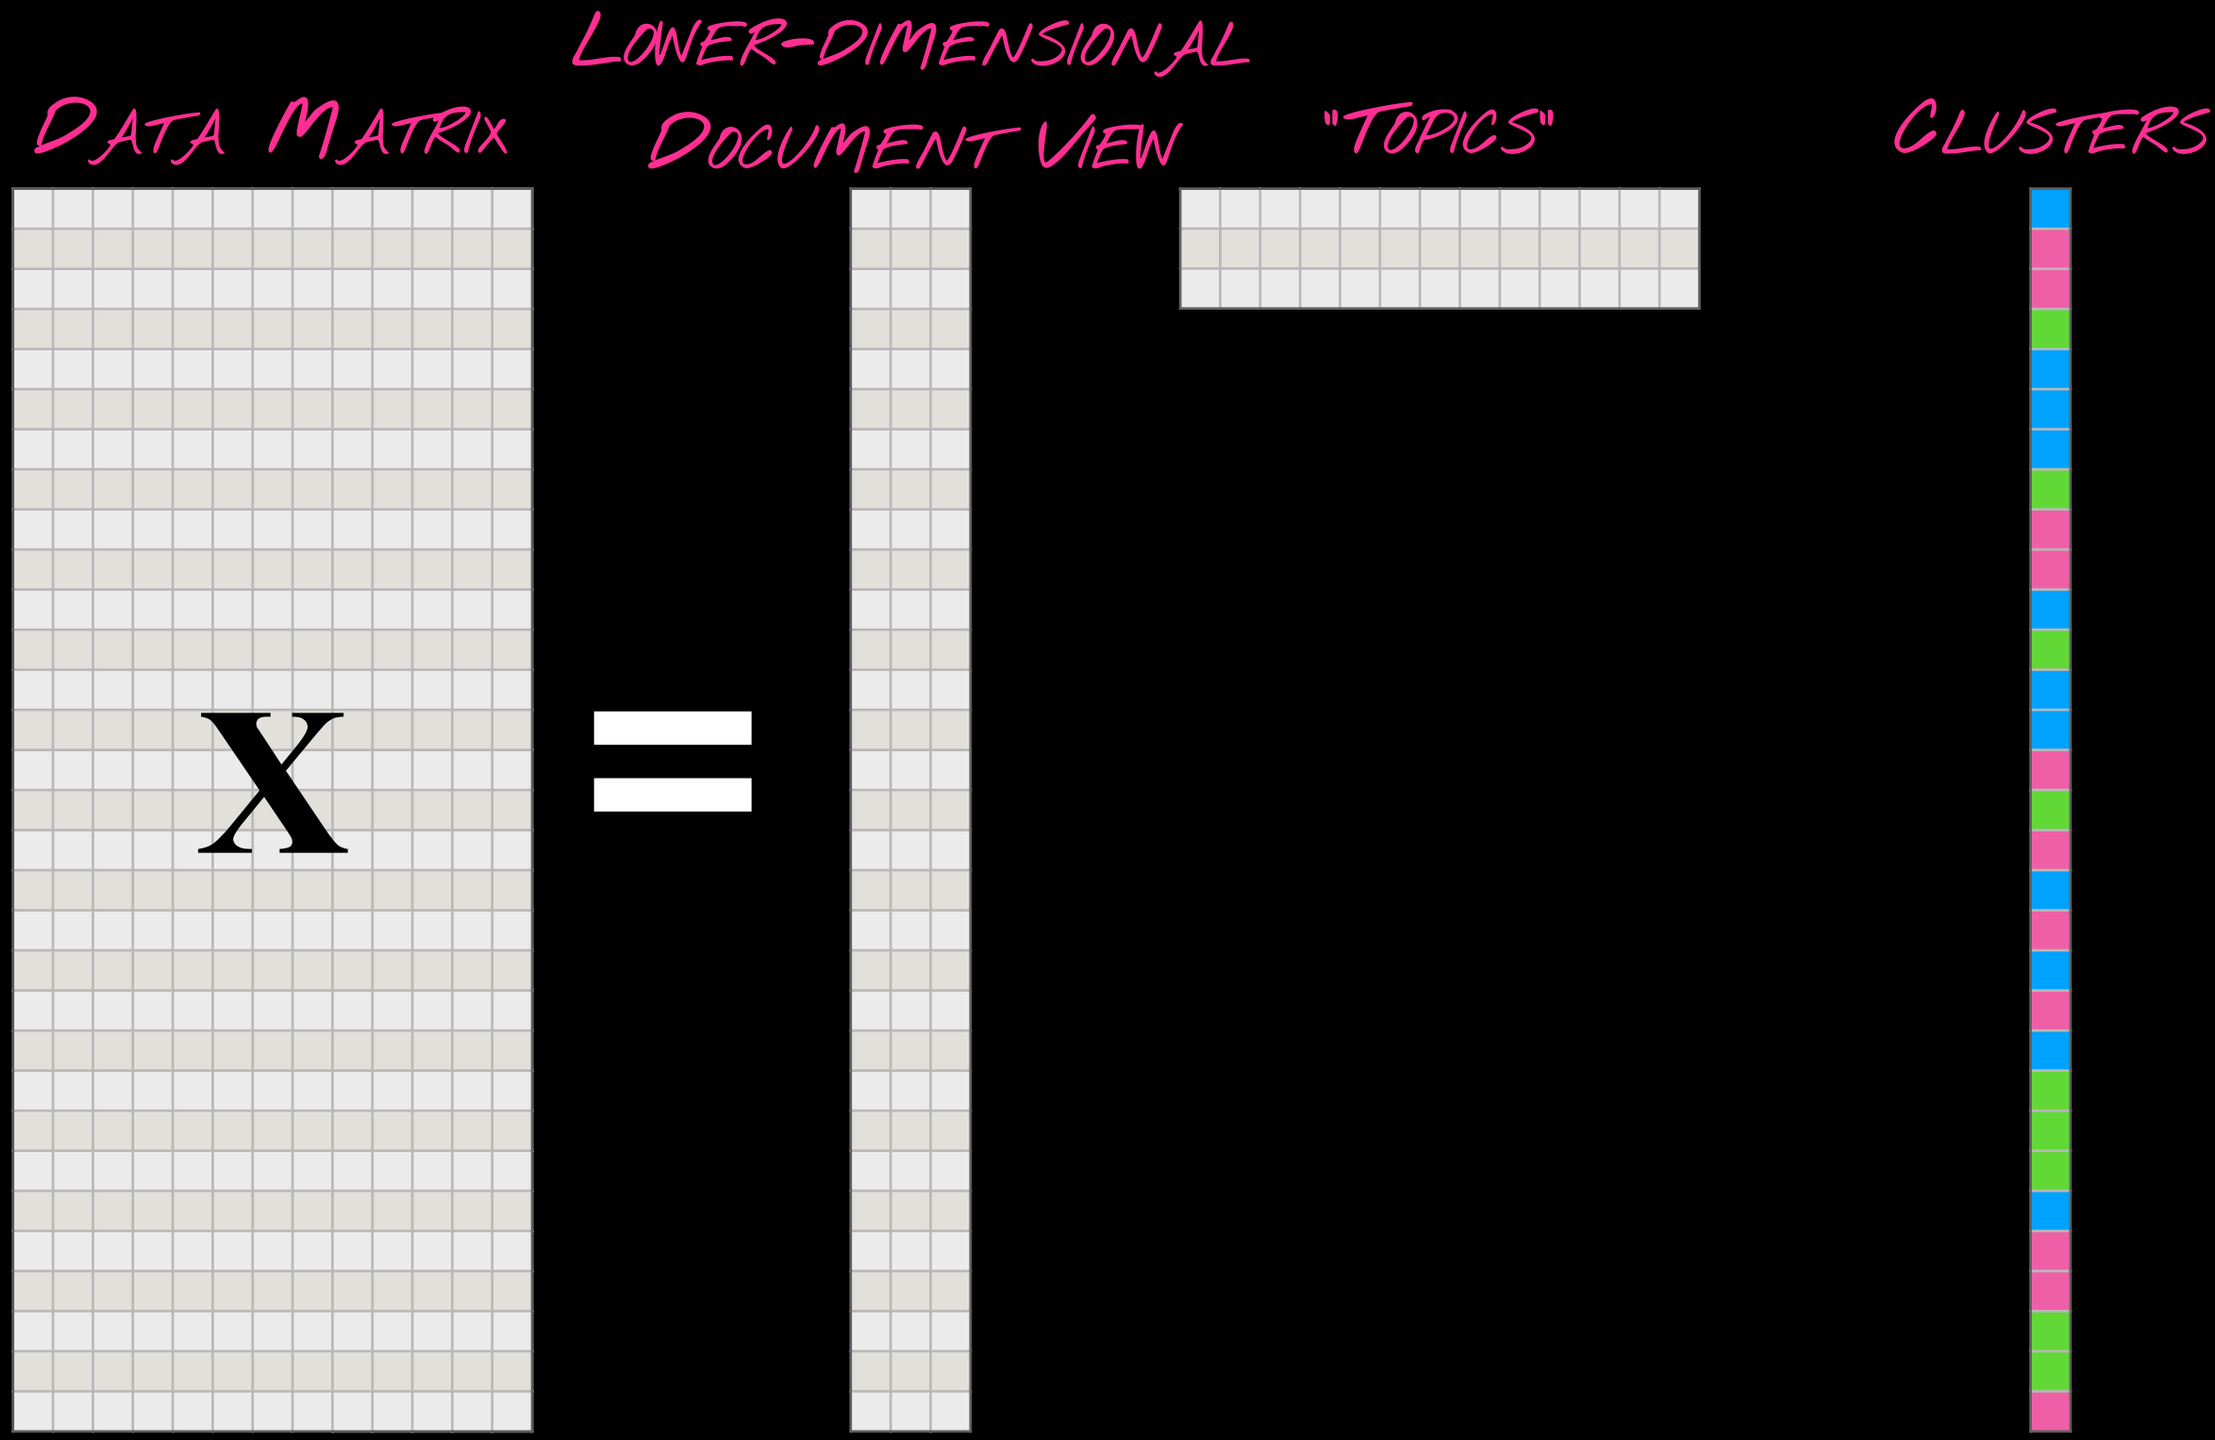

Bocconi

Costa Martins et al. (2018)

Latent Dimensions

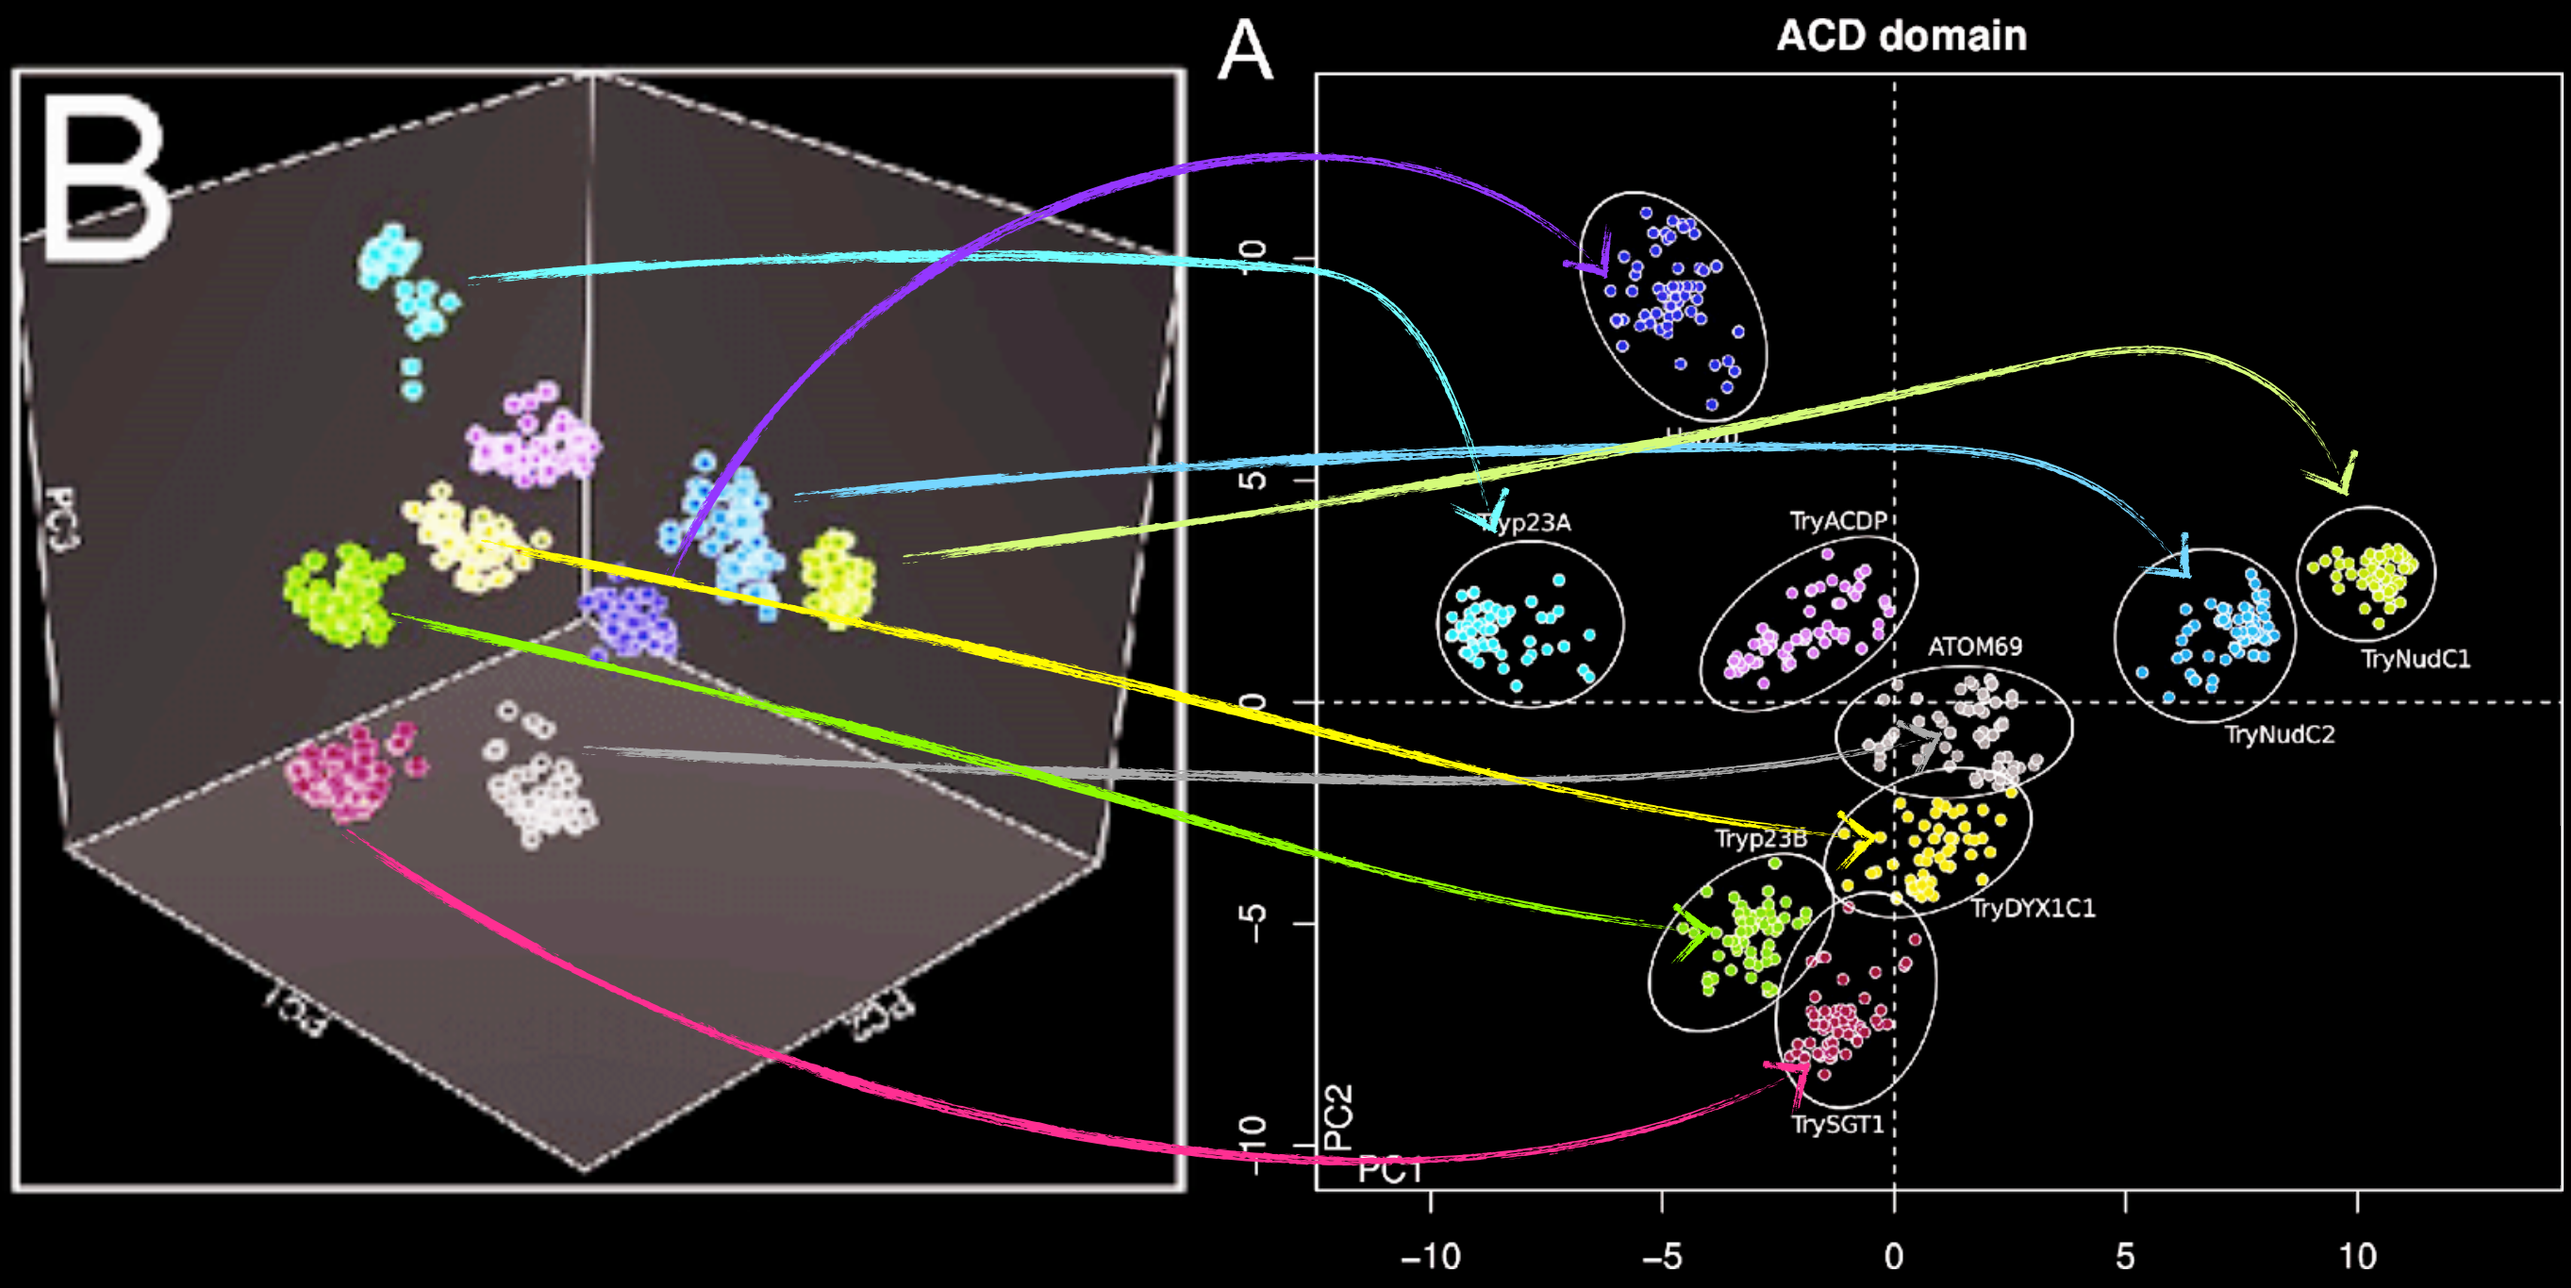

Bocconi

Goals for Today

- Learn about **matrix factorization** and its use for **semantic similarity** and **visualization**
- Learn about ***k*-means** and **agglomerative clustering**
- Learn about **evaluation** criteria

Matrix Factorization

Bocconi

Singular Value Decomposition

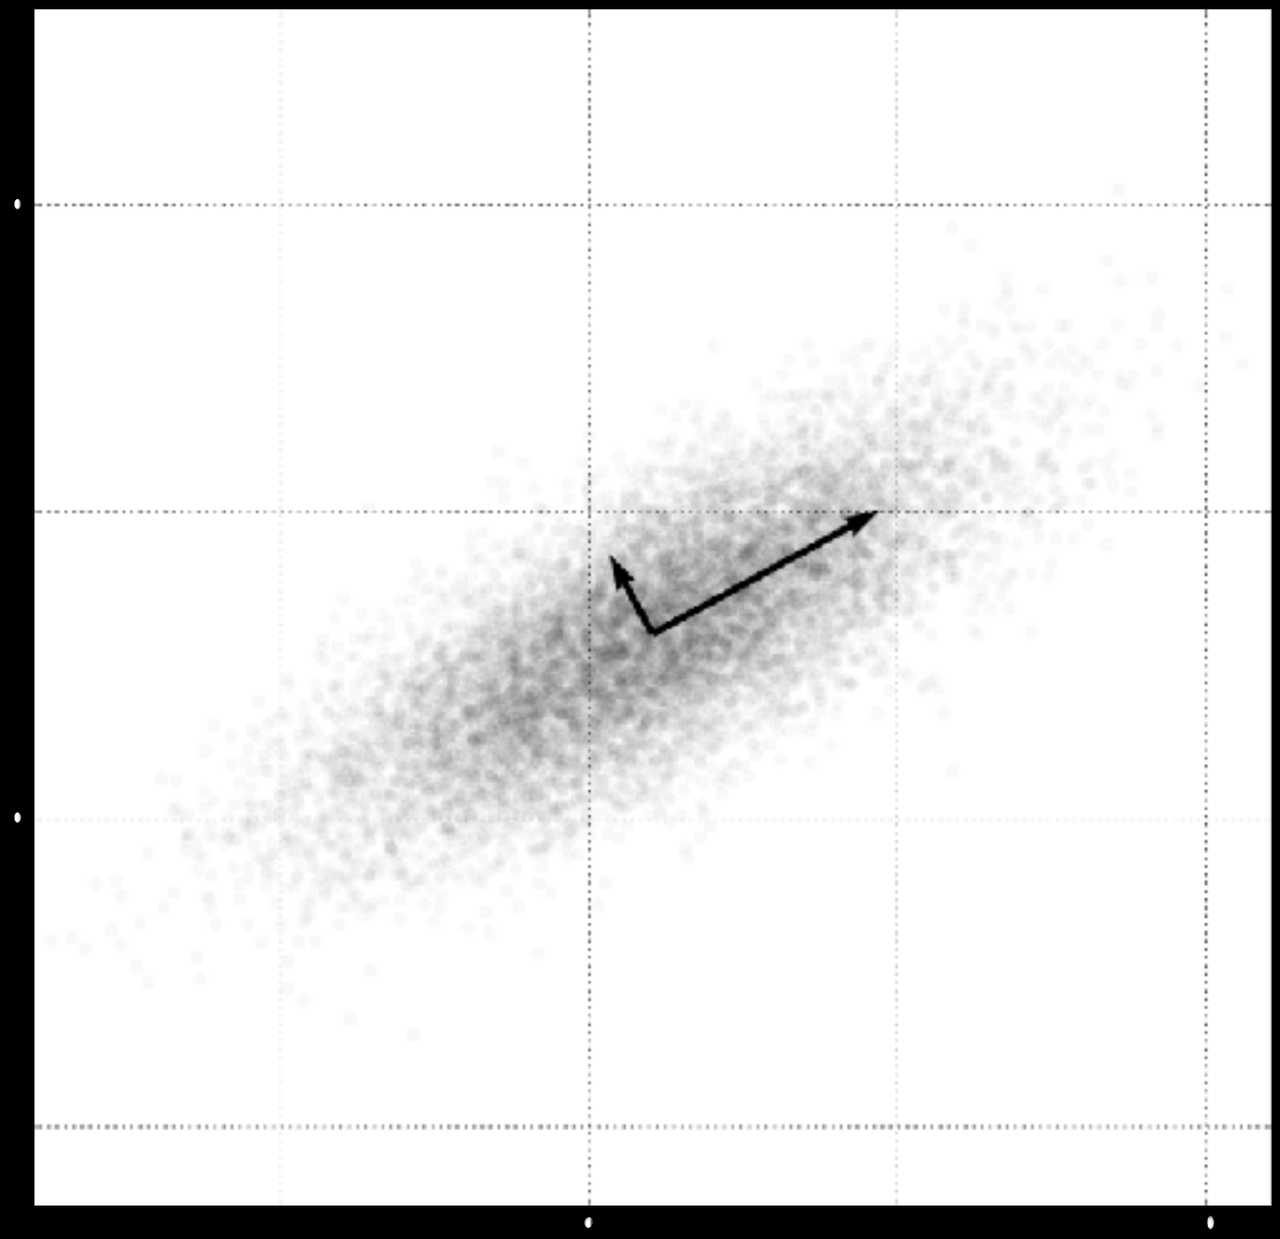

Singular Value Decomposition

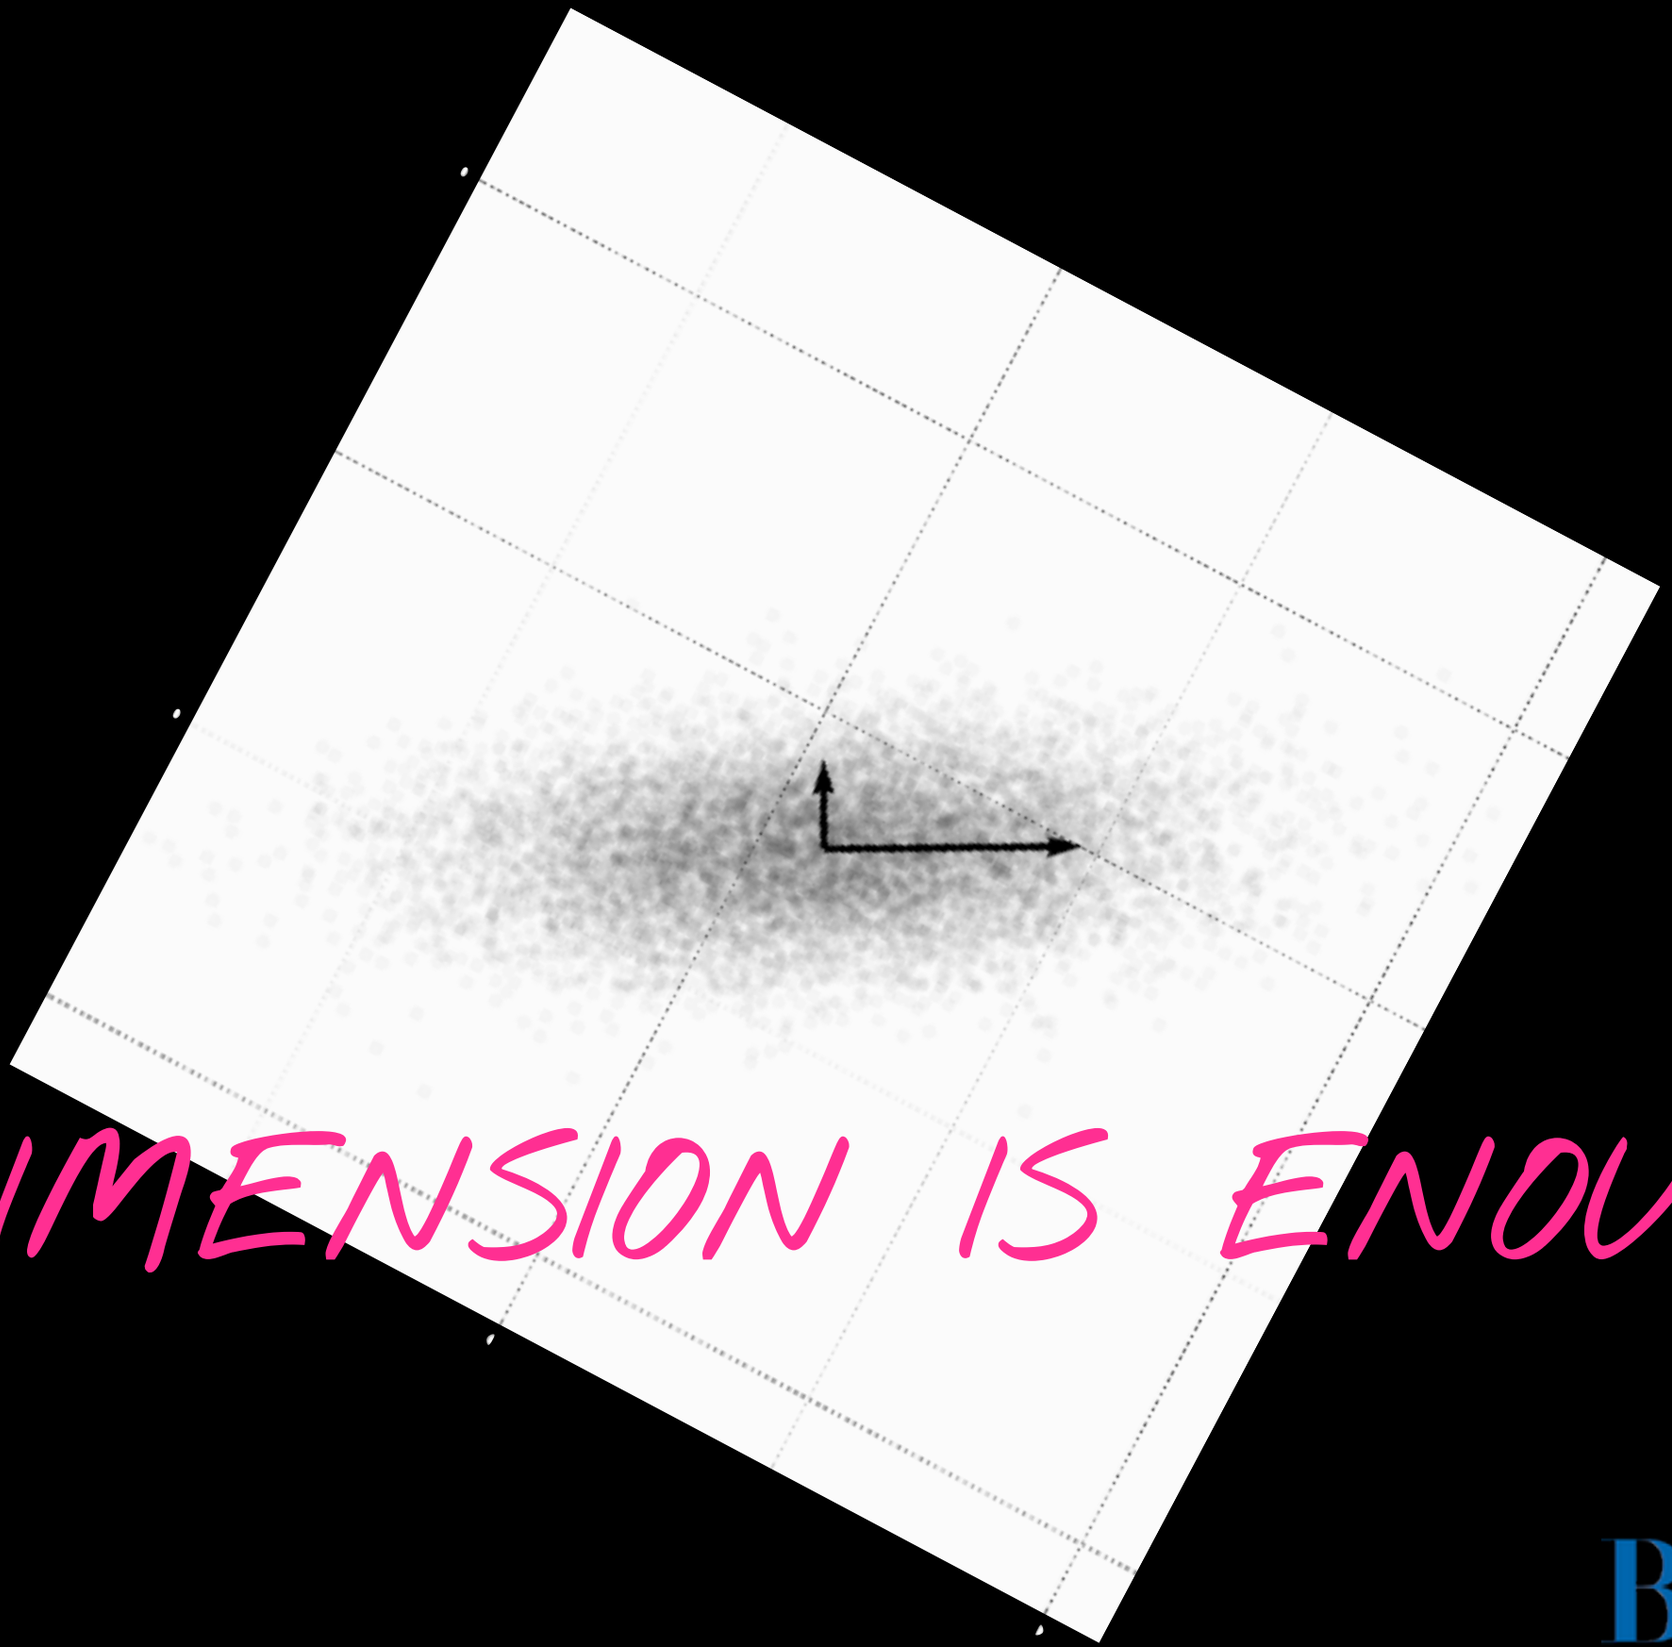

I DIMENSION IS ENOUGH!

Singular Value Decomposition

- “principal component analysis”: discover the dimensions that matter
- idea: matrix is made up of few hidden dimensions
- Dimensions correspond to **documents**, **terms**, and latent **concepts**

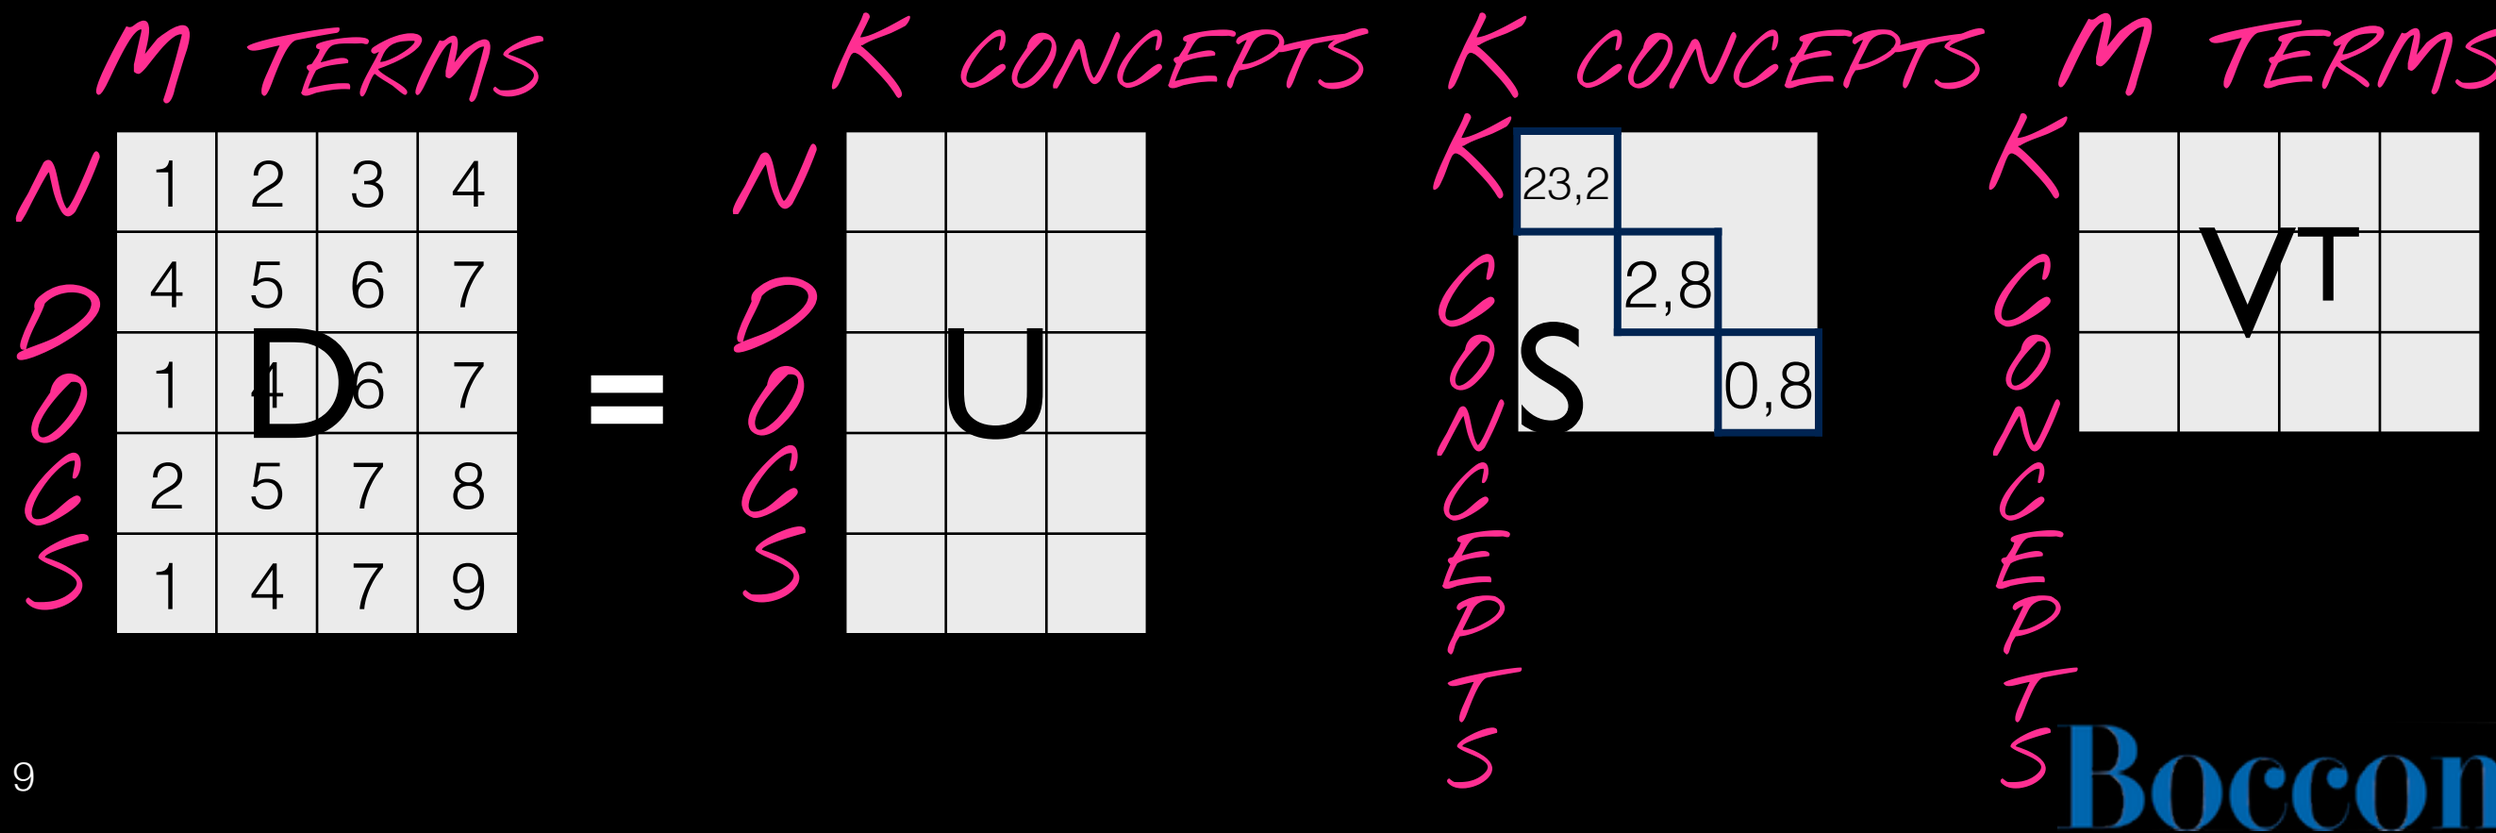

Bocconi

Singular Value Decomposition

- reduce principal components/concepts to smaller number

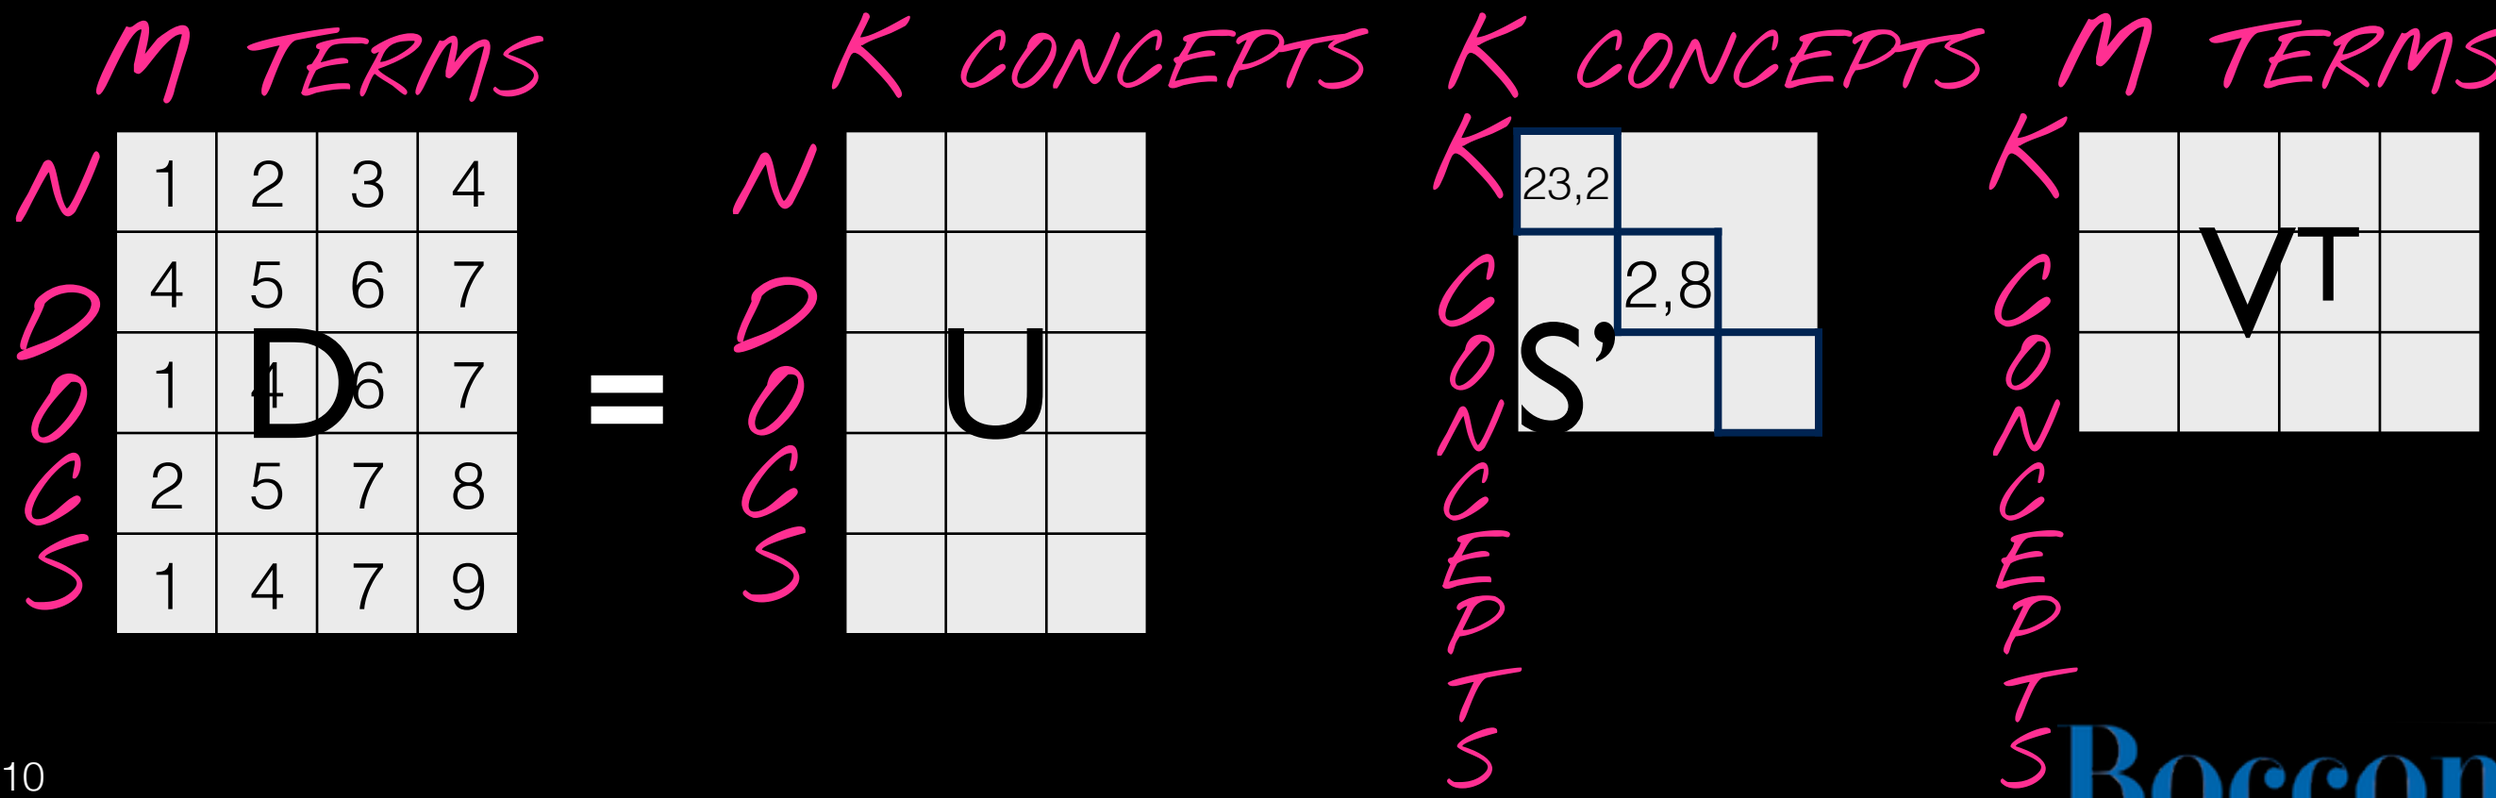

Singular Value Decomposition

- reconstruct original matrix in new concept space:
  Latent Semantic Analysis

Bocconi

Non-negative Matrix Factorization

- Use only positive values
- Find approximation of two components

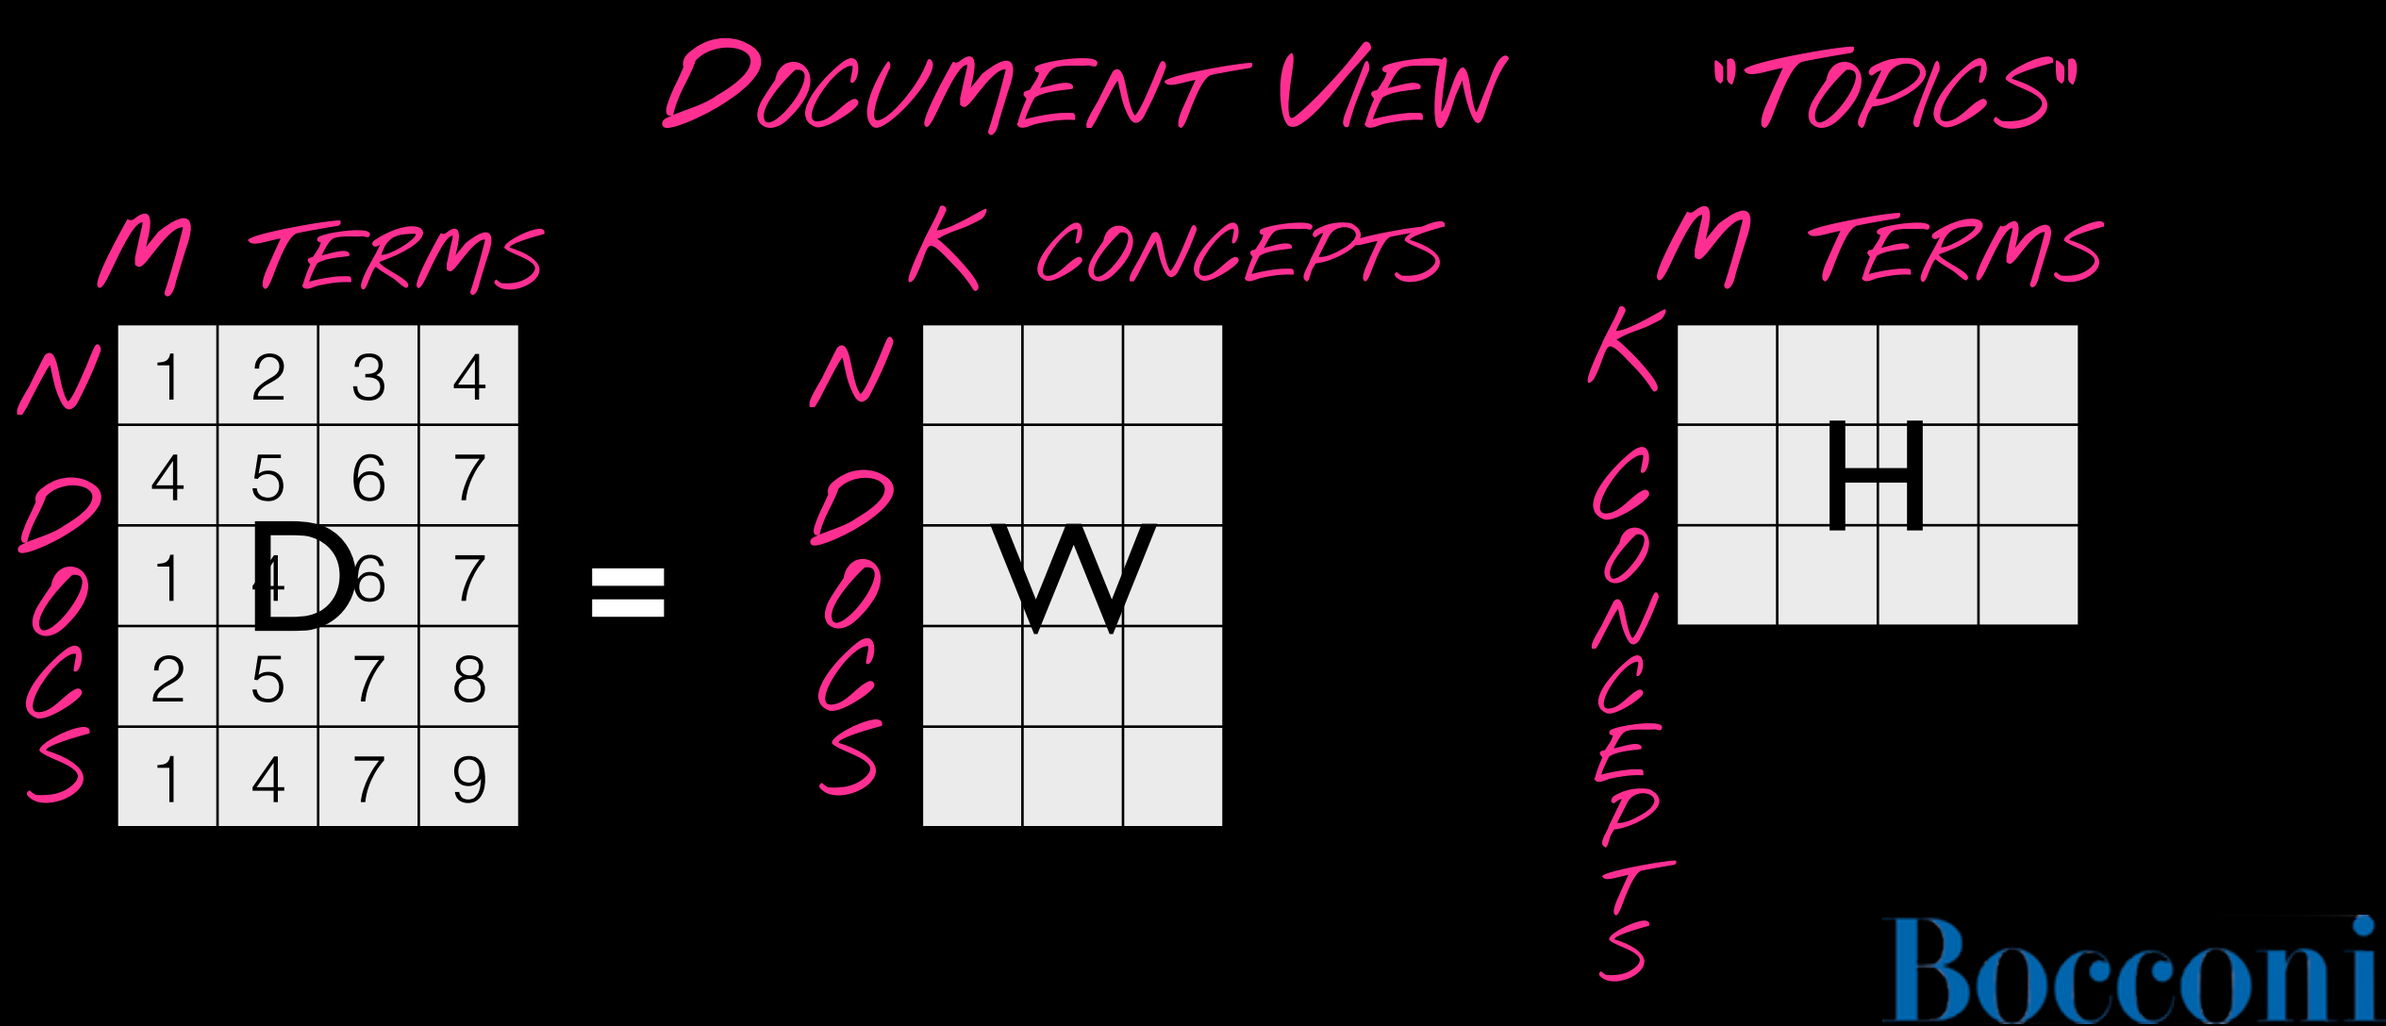

Bocconi

Comparison

	SVD	NMF
Negative values (embeddings) as input?	yes	no
#components	3: U, S, V	2: W, H
document view?	yes: U	yes: W
term view?	yes: V	yes: H
strength ranking?	yes: S	no
exact?	yes	no
"topic" quality	mixed	better
sparsity	low	medium

Yes, but: What is it Good for?

- Find latent **topic** dimensions (alternative: LDA)
- Find **word similarity** in latent space (alternative: Word2Vec)
- Find **document similarity** in latent space (alternative: Doc2Vec)
- Reduce dimensionality for **visualization**

Latent Word Dimension Topics

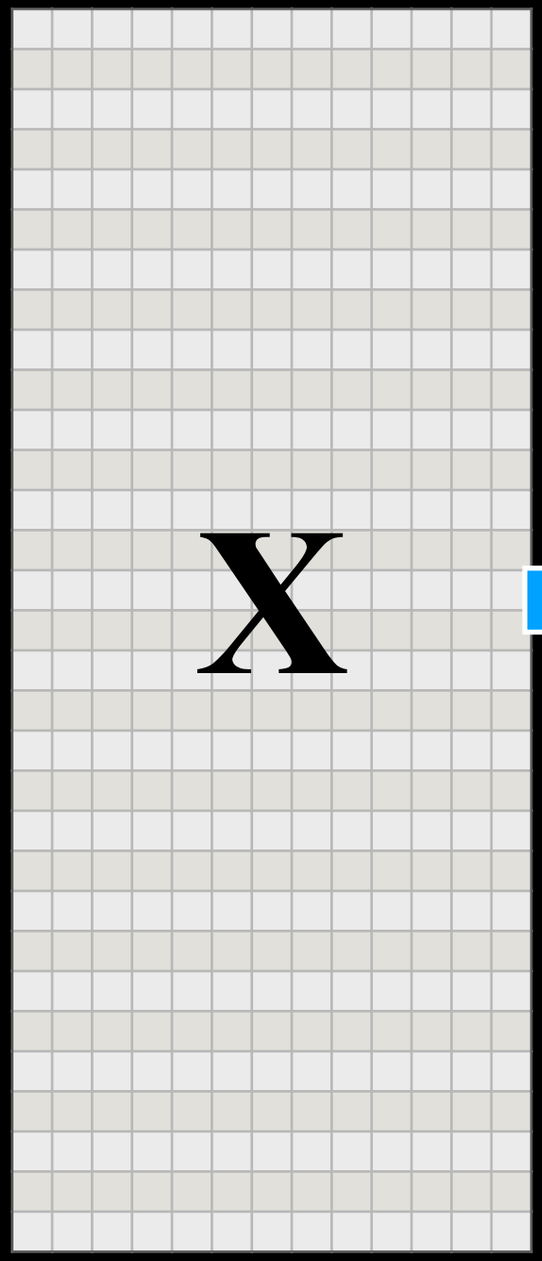

NMF(10) FOR MOBY DICK

AHAB	ahab, captain, cried, captain ahab, cried ahab
STRUCTURE	chapter, folio, octavo, ii, iii
???	like, ship, sea, time, way
MEN	man, old, old man, look, young man
???	oh, life, starbuck, sweet, god
CHARACTERS	said, stubb, queequeg, don, starbuck
???	sir, aye, let, shall, think
OLD-TIMEY	thou, thee, thy, st, god
WHALES	whale, sperm, sperm whale, white, white whale
???	ye, look, say, ye ye, men

Dimensionality Reduction for Visualizations

Bocconi

Hovy & Purschke, 2018

Dimensions as RGB

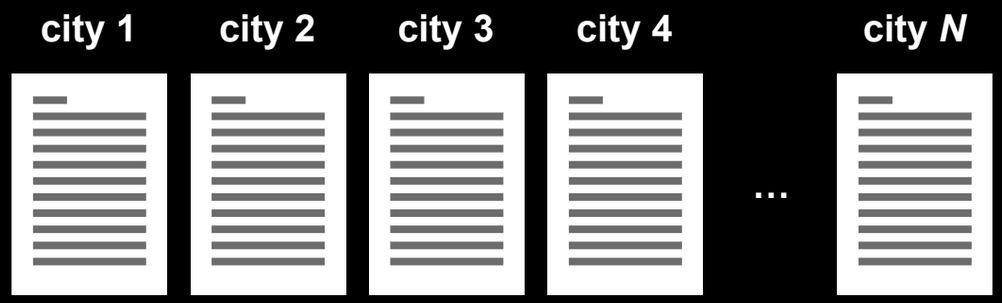

Doc2Vec

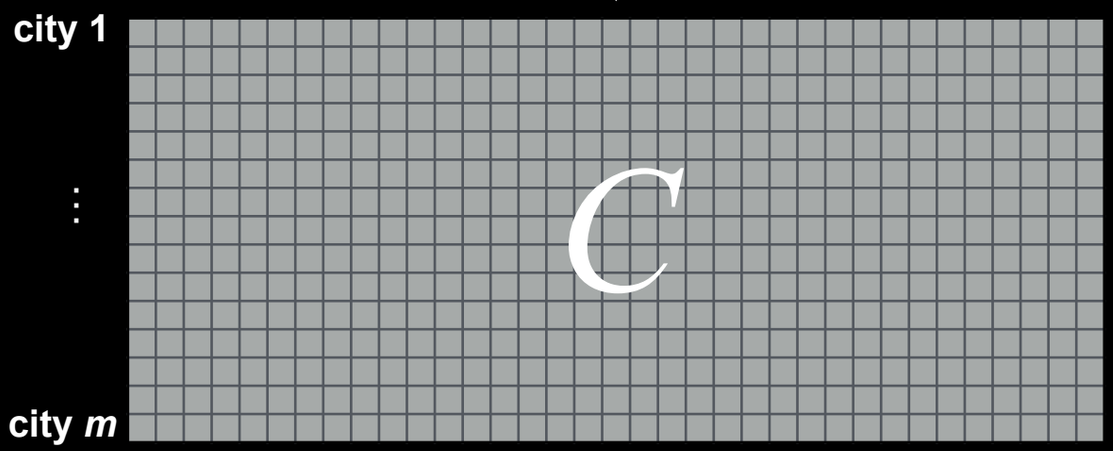

=

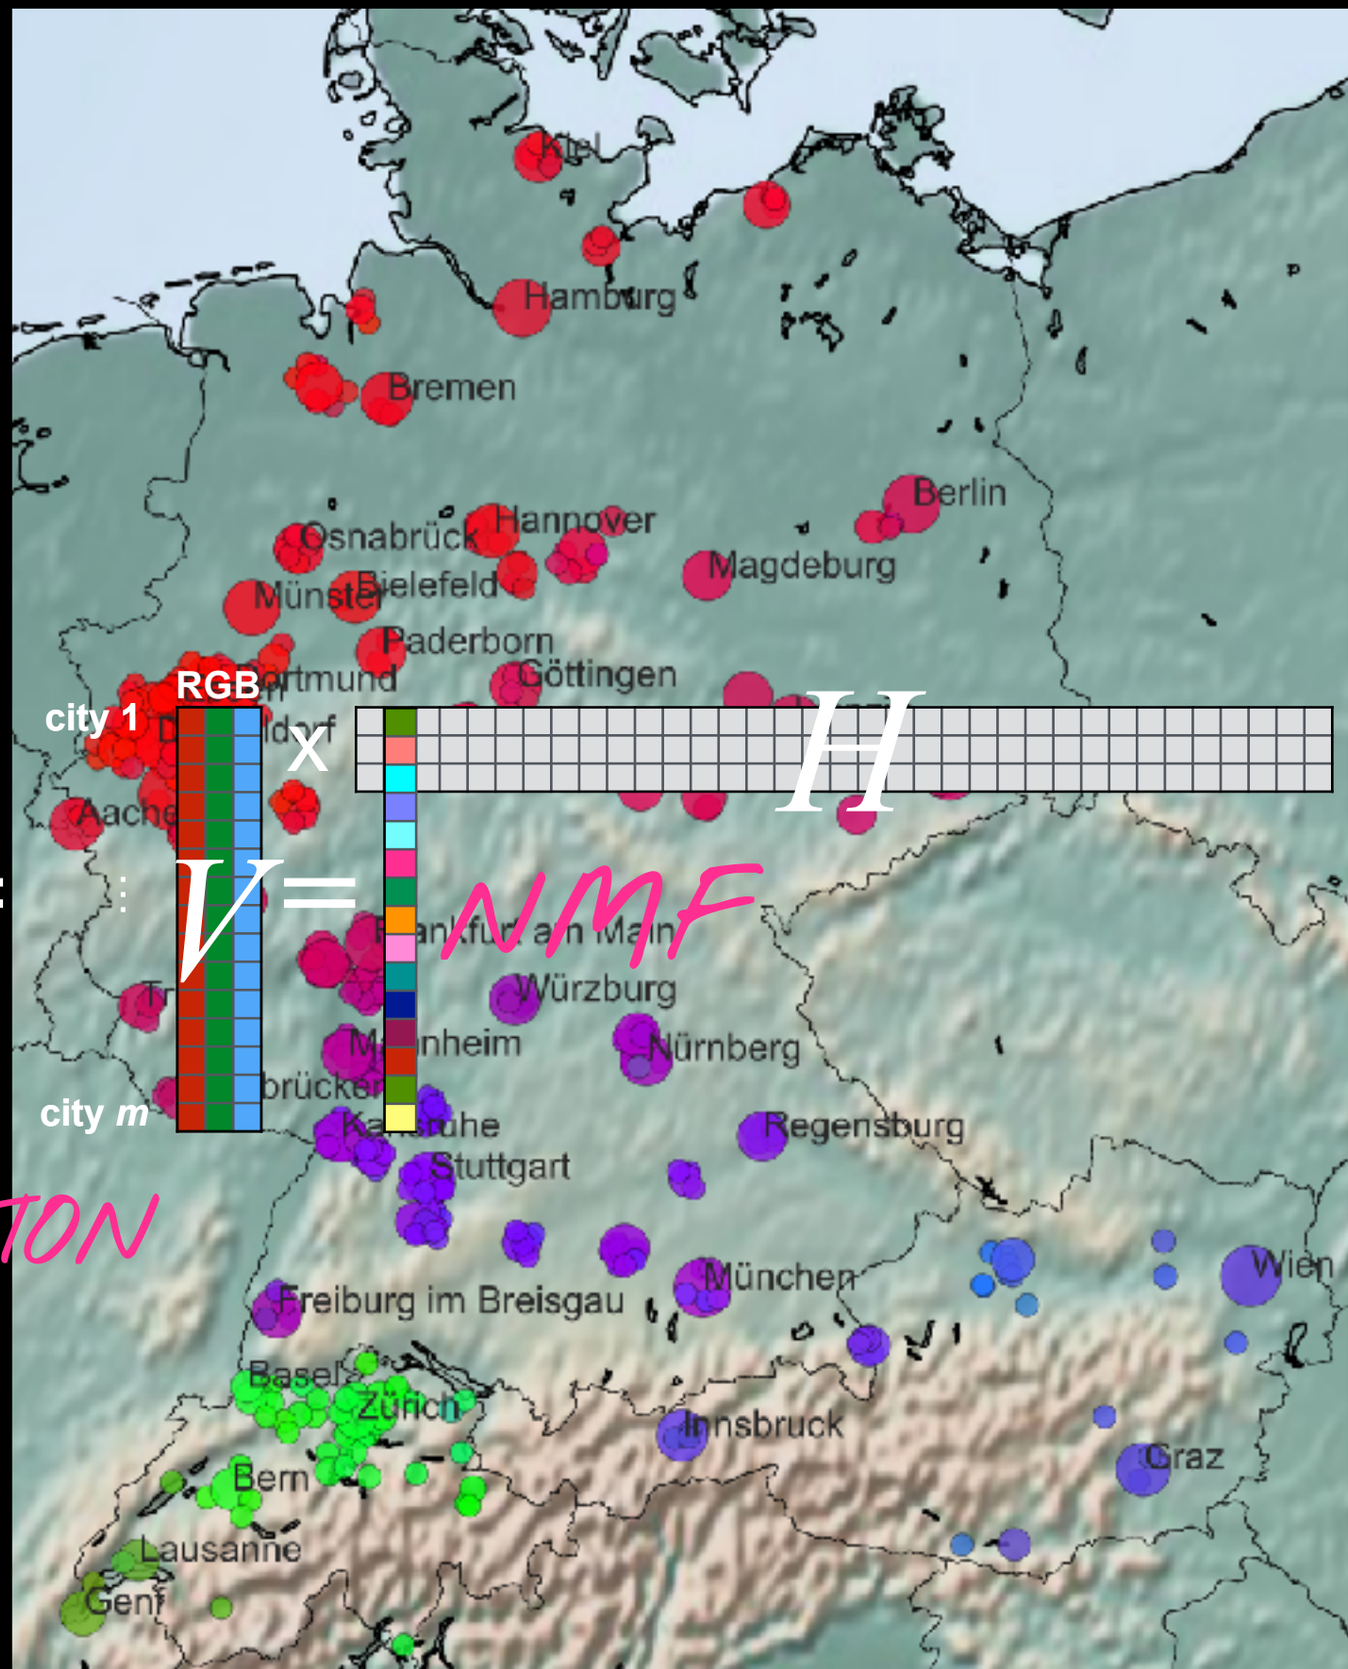

Dense Representation

Hovy et al., 2019

Dimensions as RGB

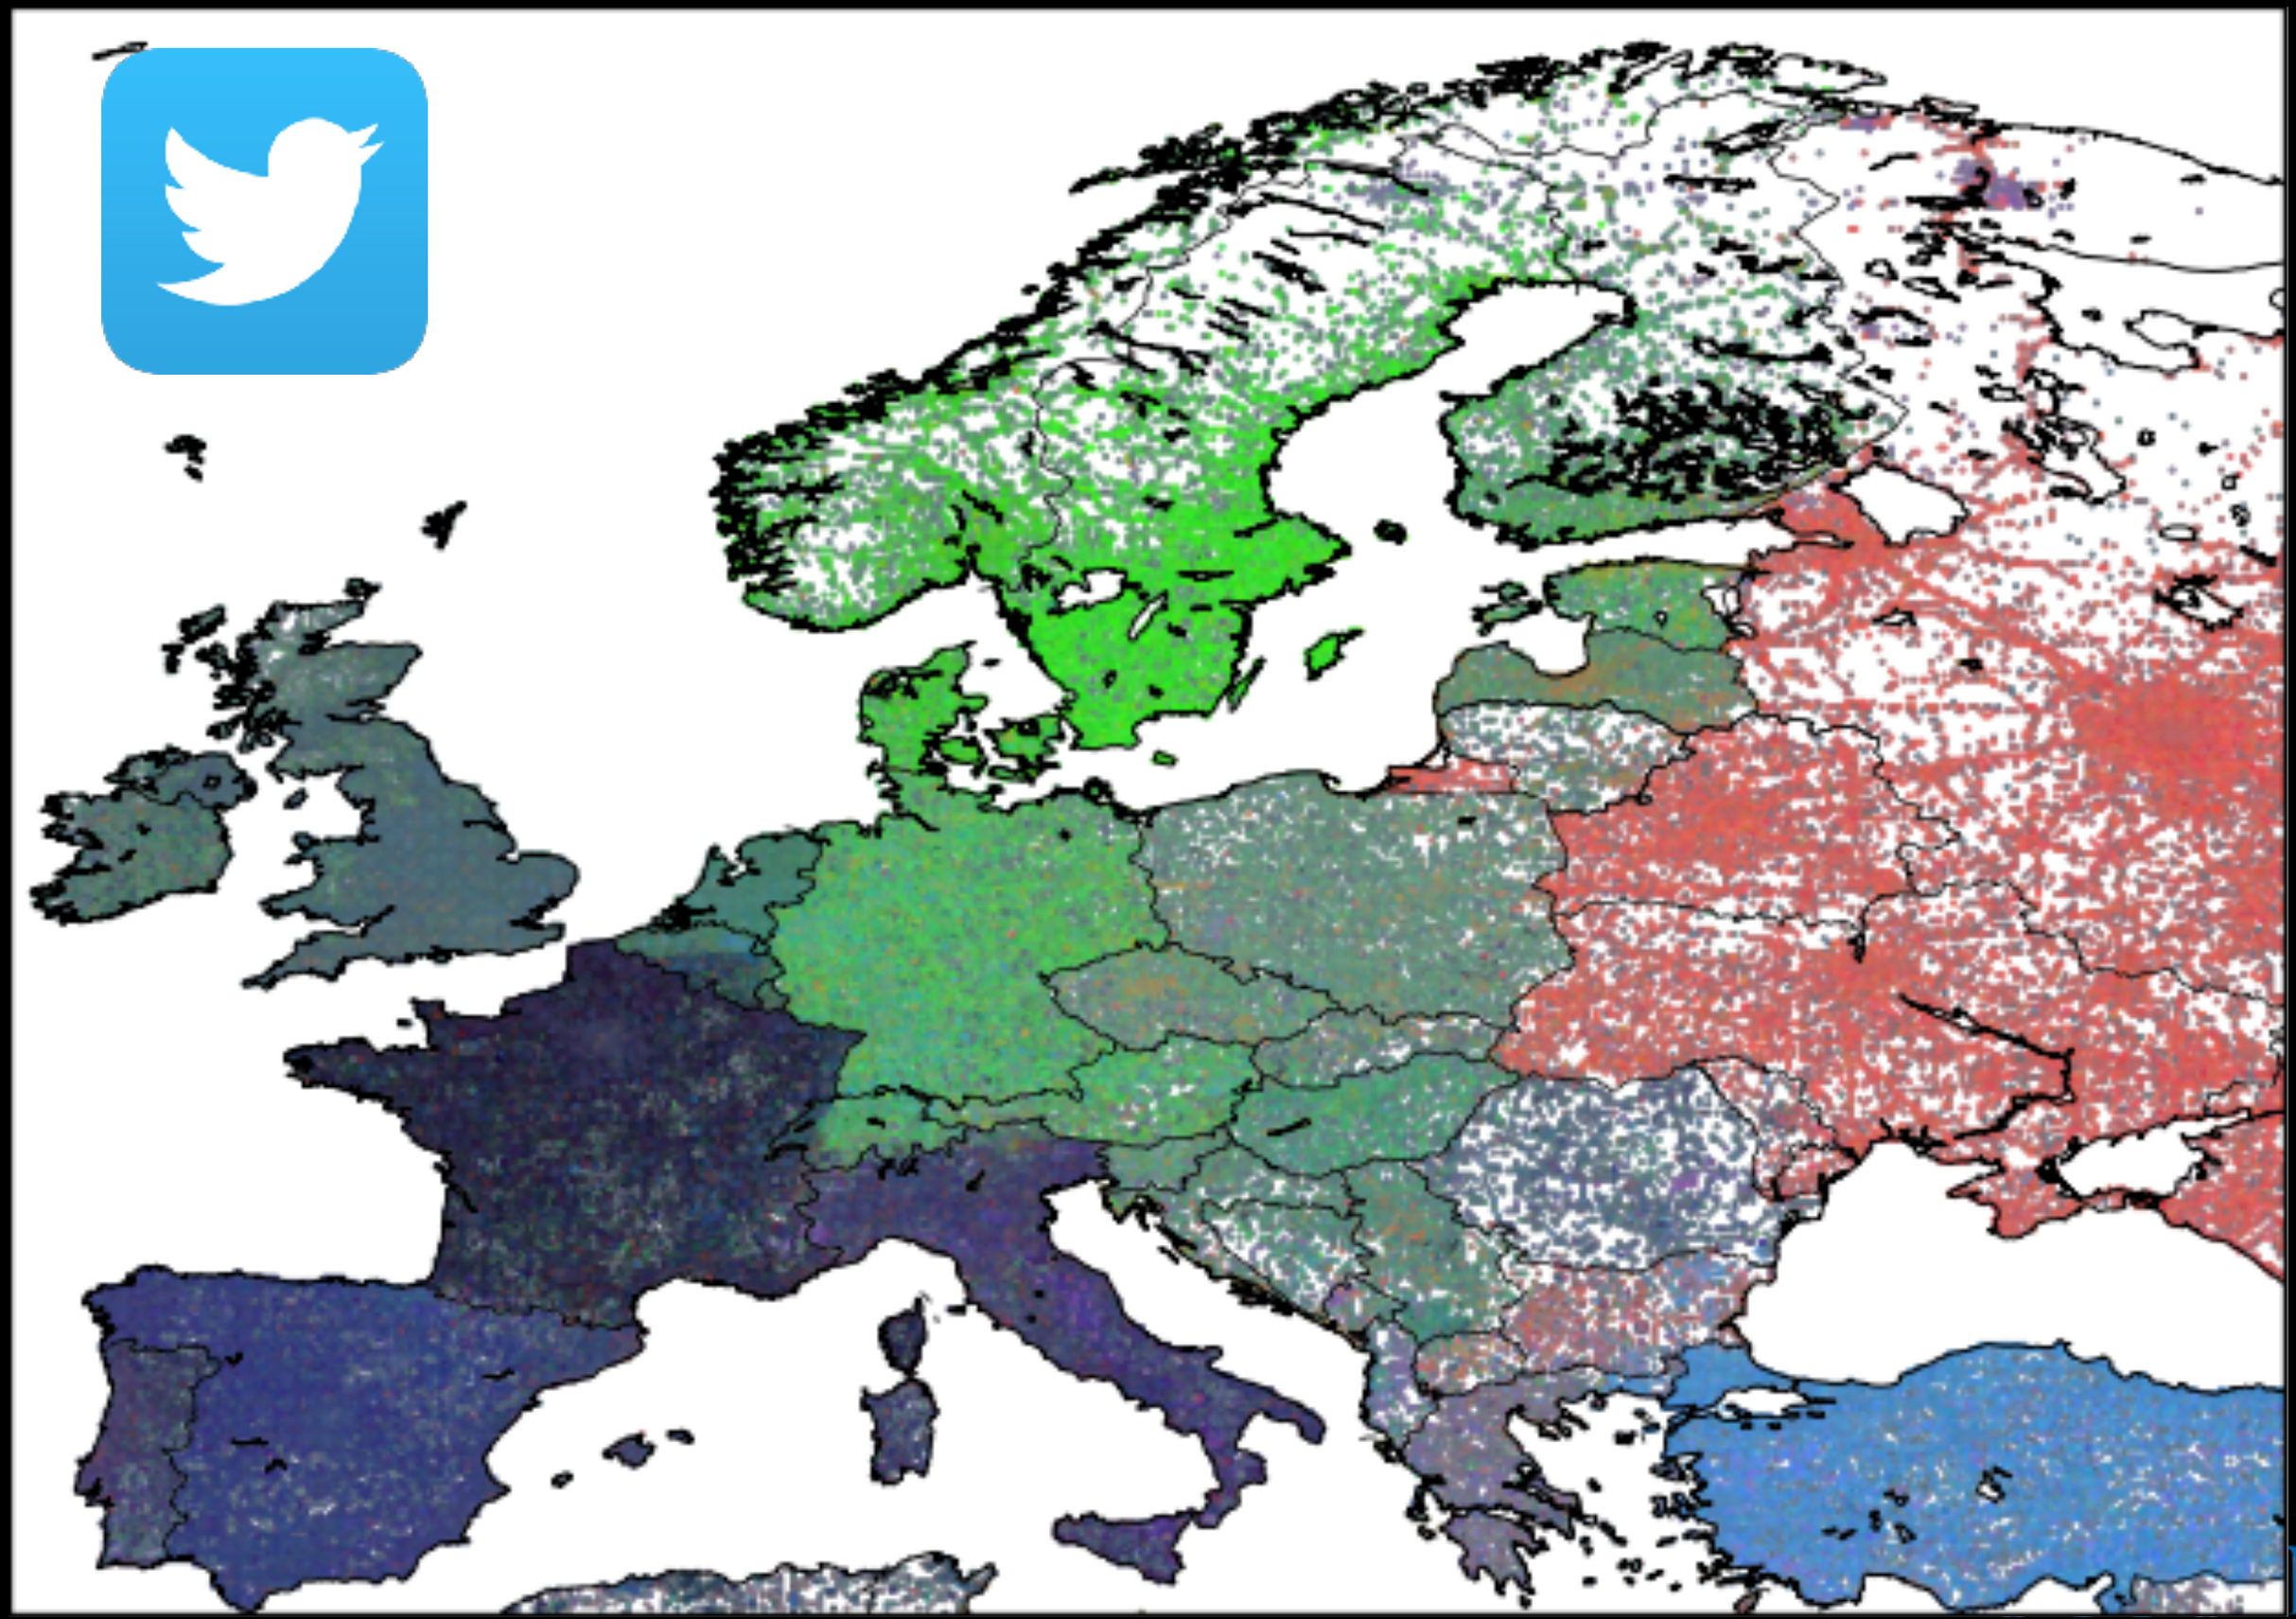

Dimensions as Coordinates

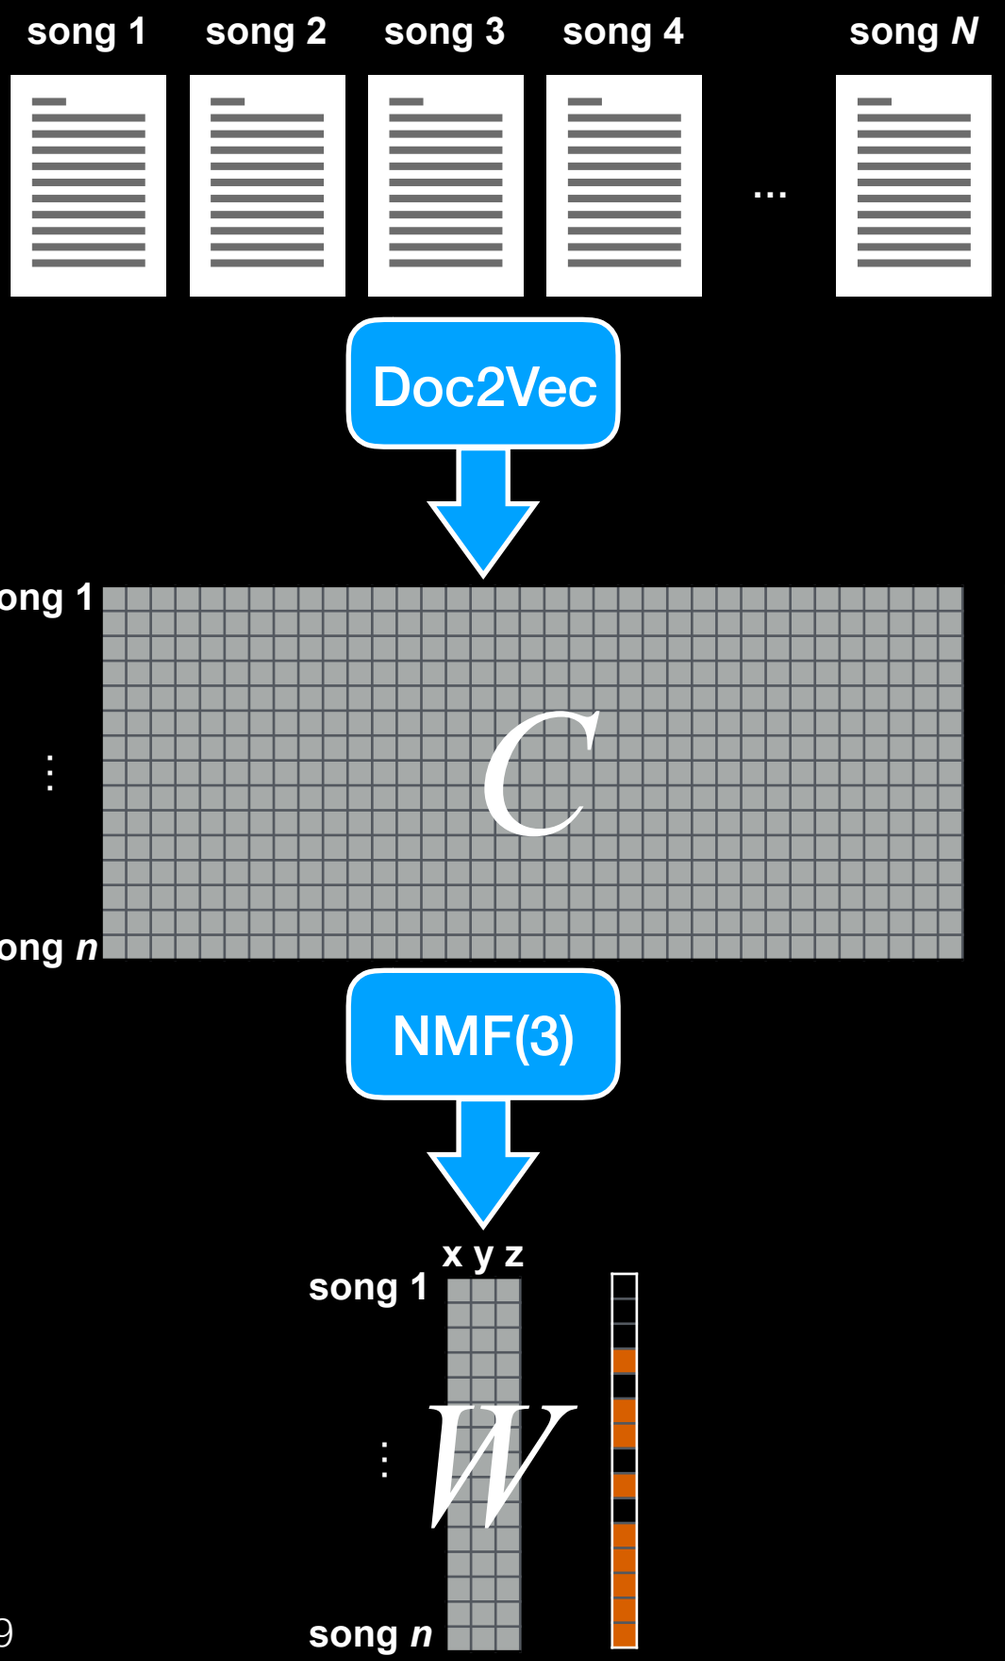

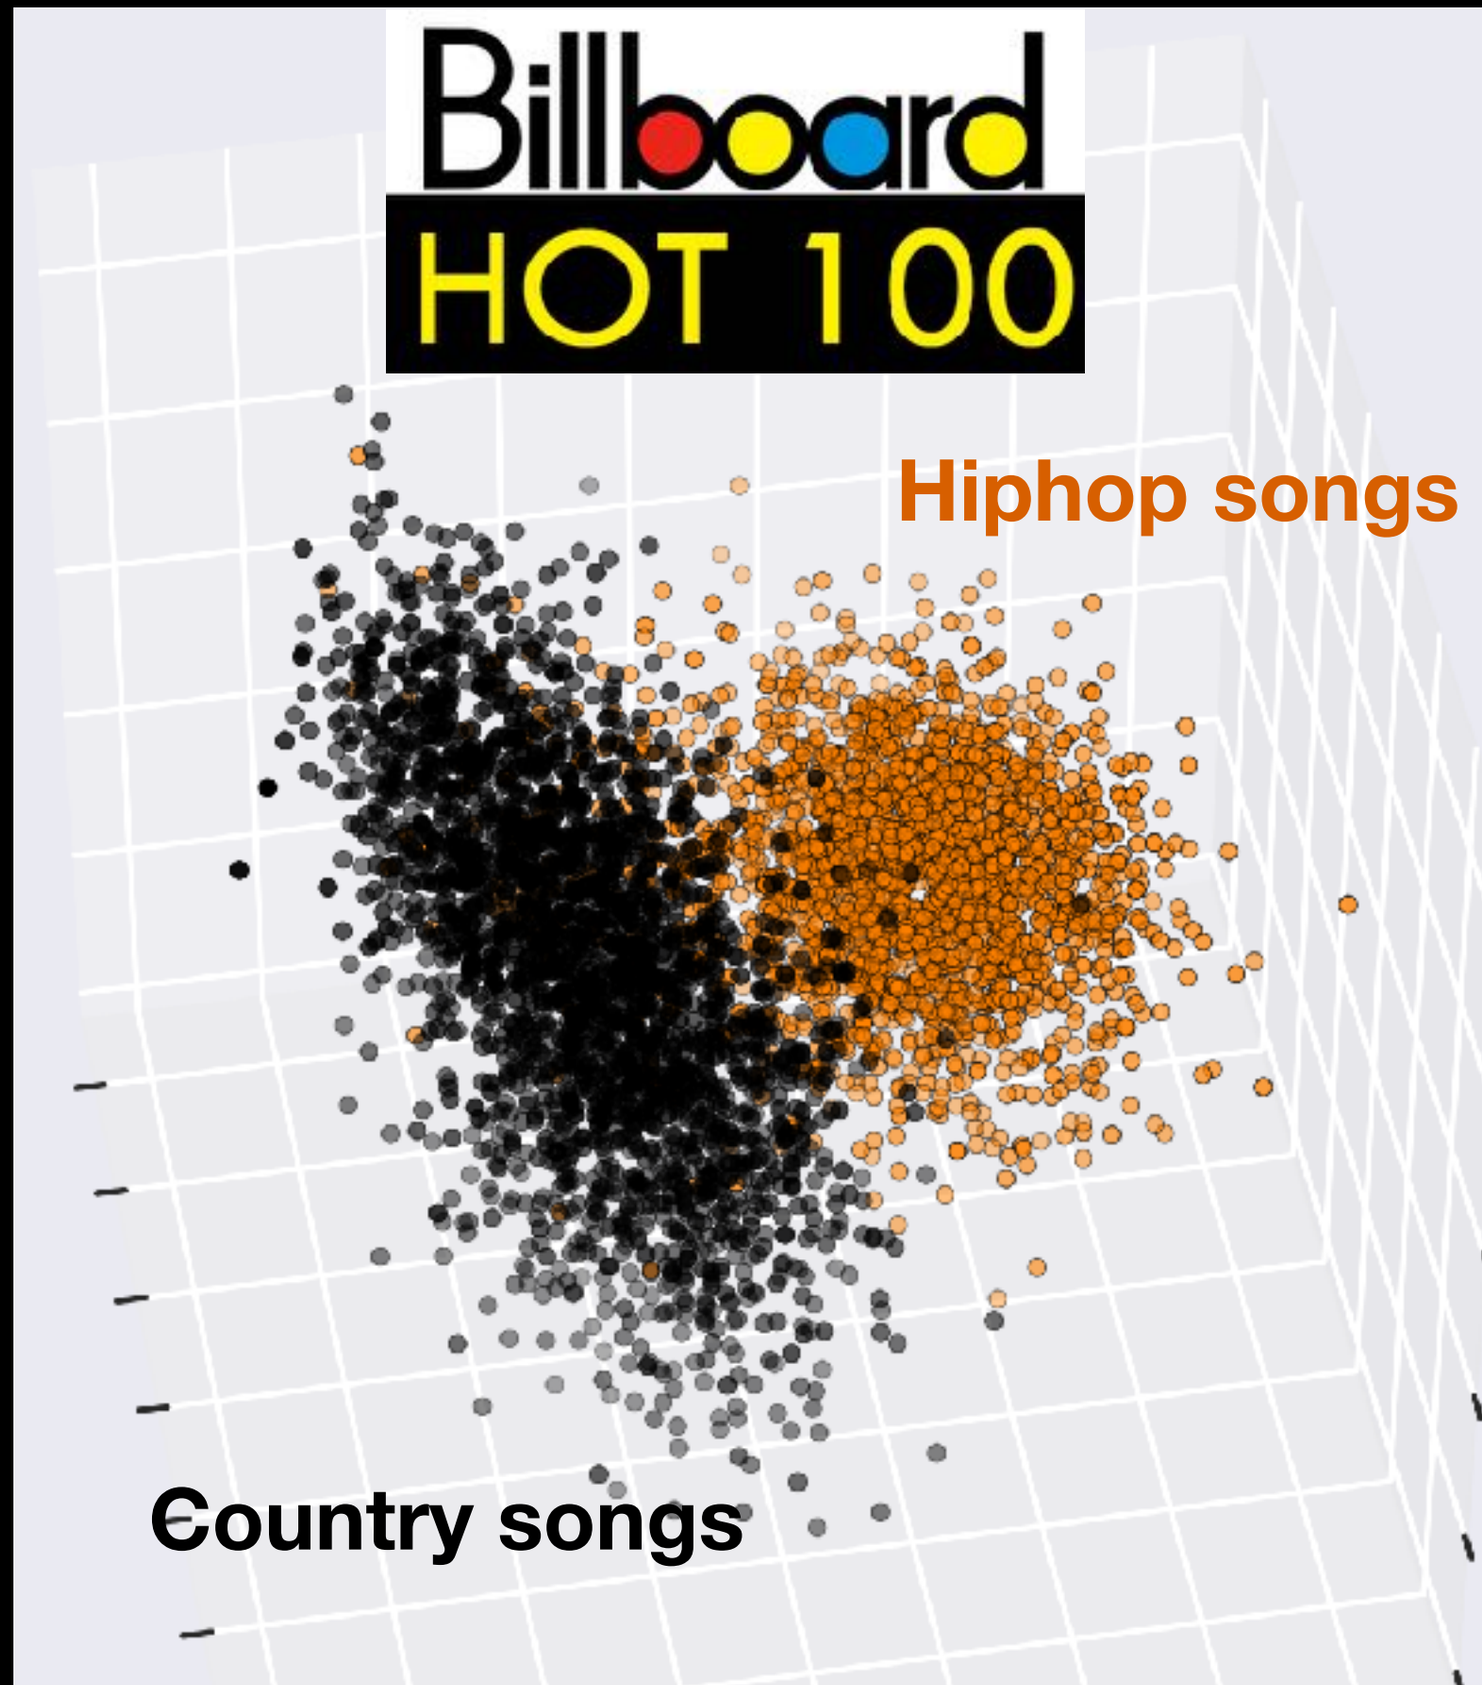

Van der Maaten & Hinton, 2008

t-SNE

- Map/preserve neighborhood structure of high-dimensional space in lower-dimensional space

- Minimize difference between probability distributions over neighbors in both dimensions

$P(n_j|n_i)$

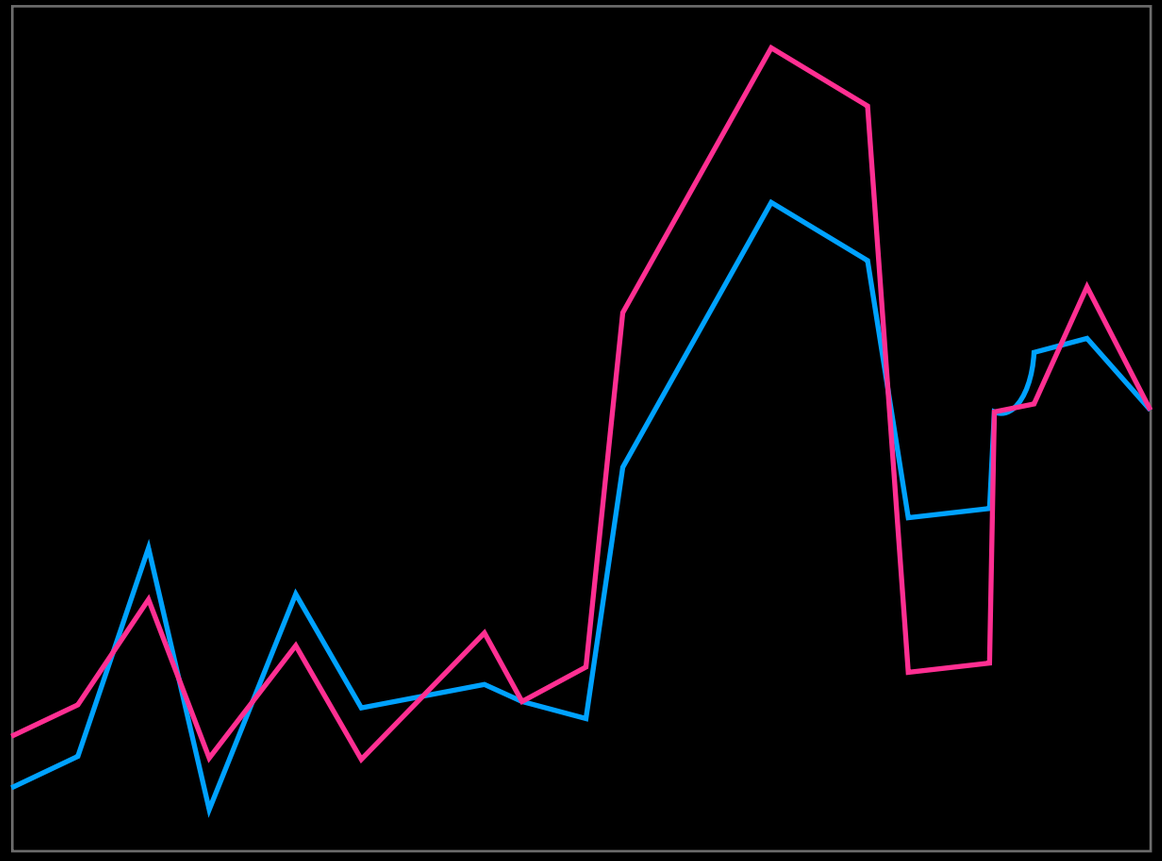

+ Nicer spatial images

– Stochastic (many runs)

– Many parameters

t-SNE

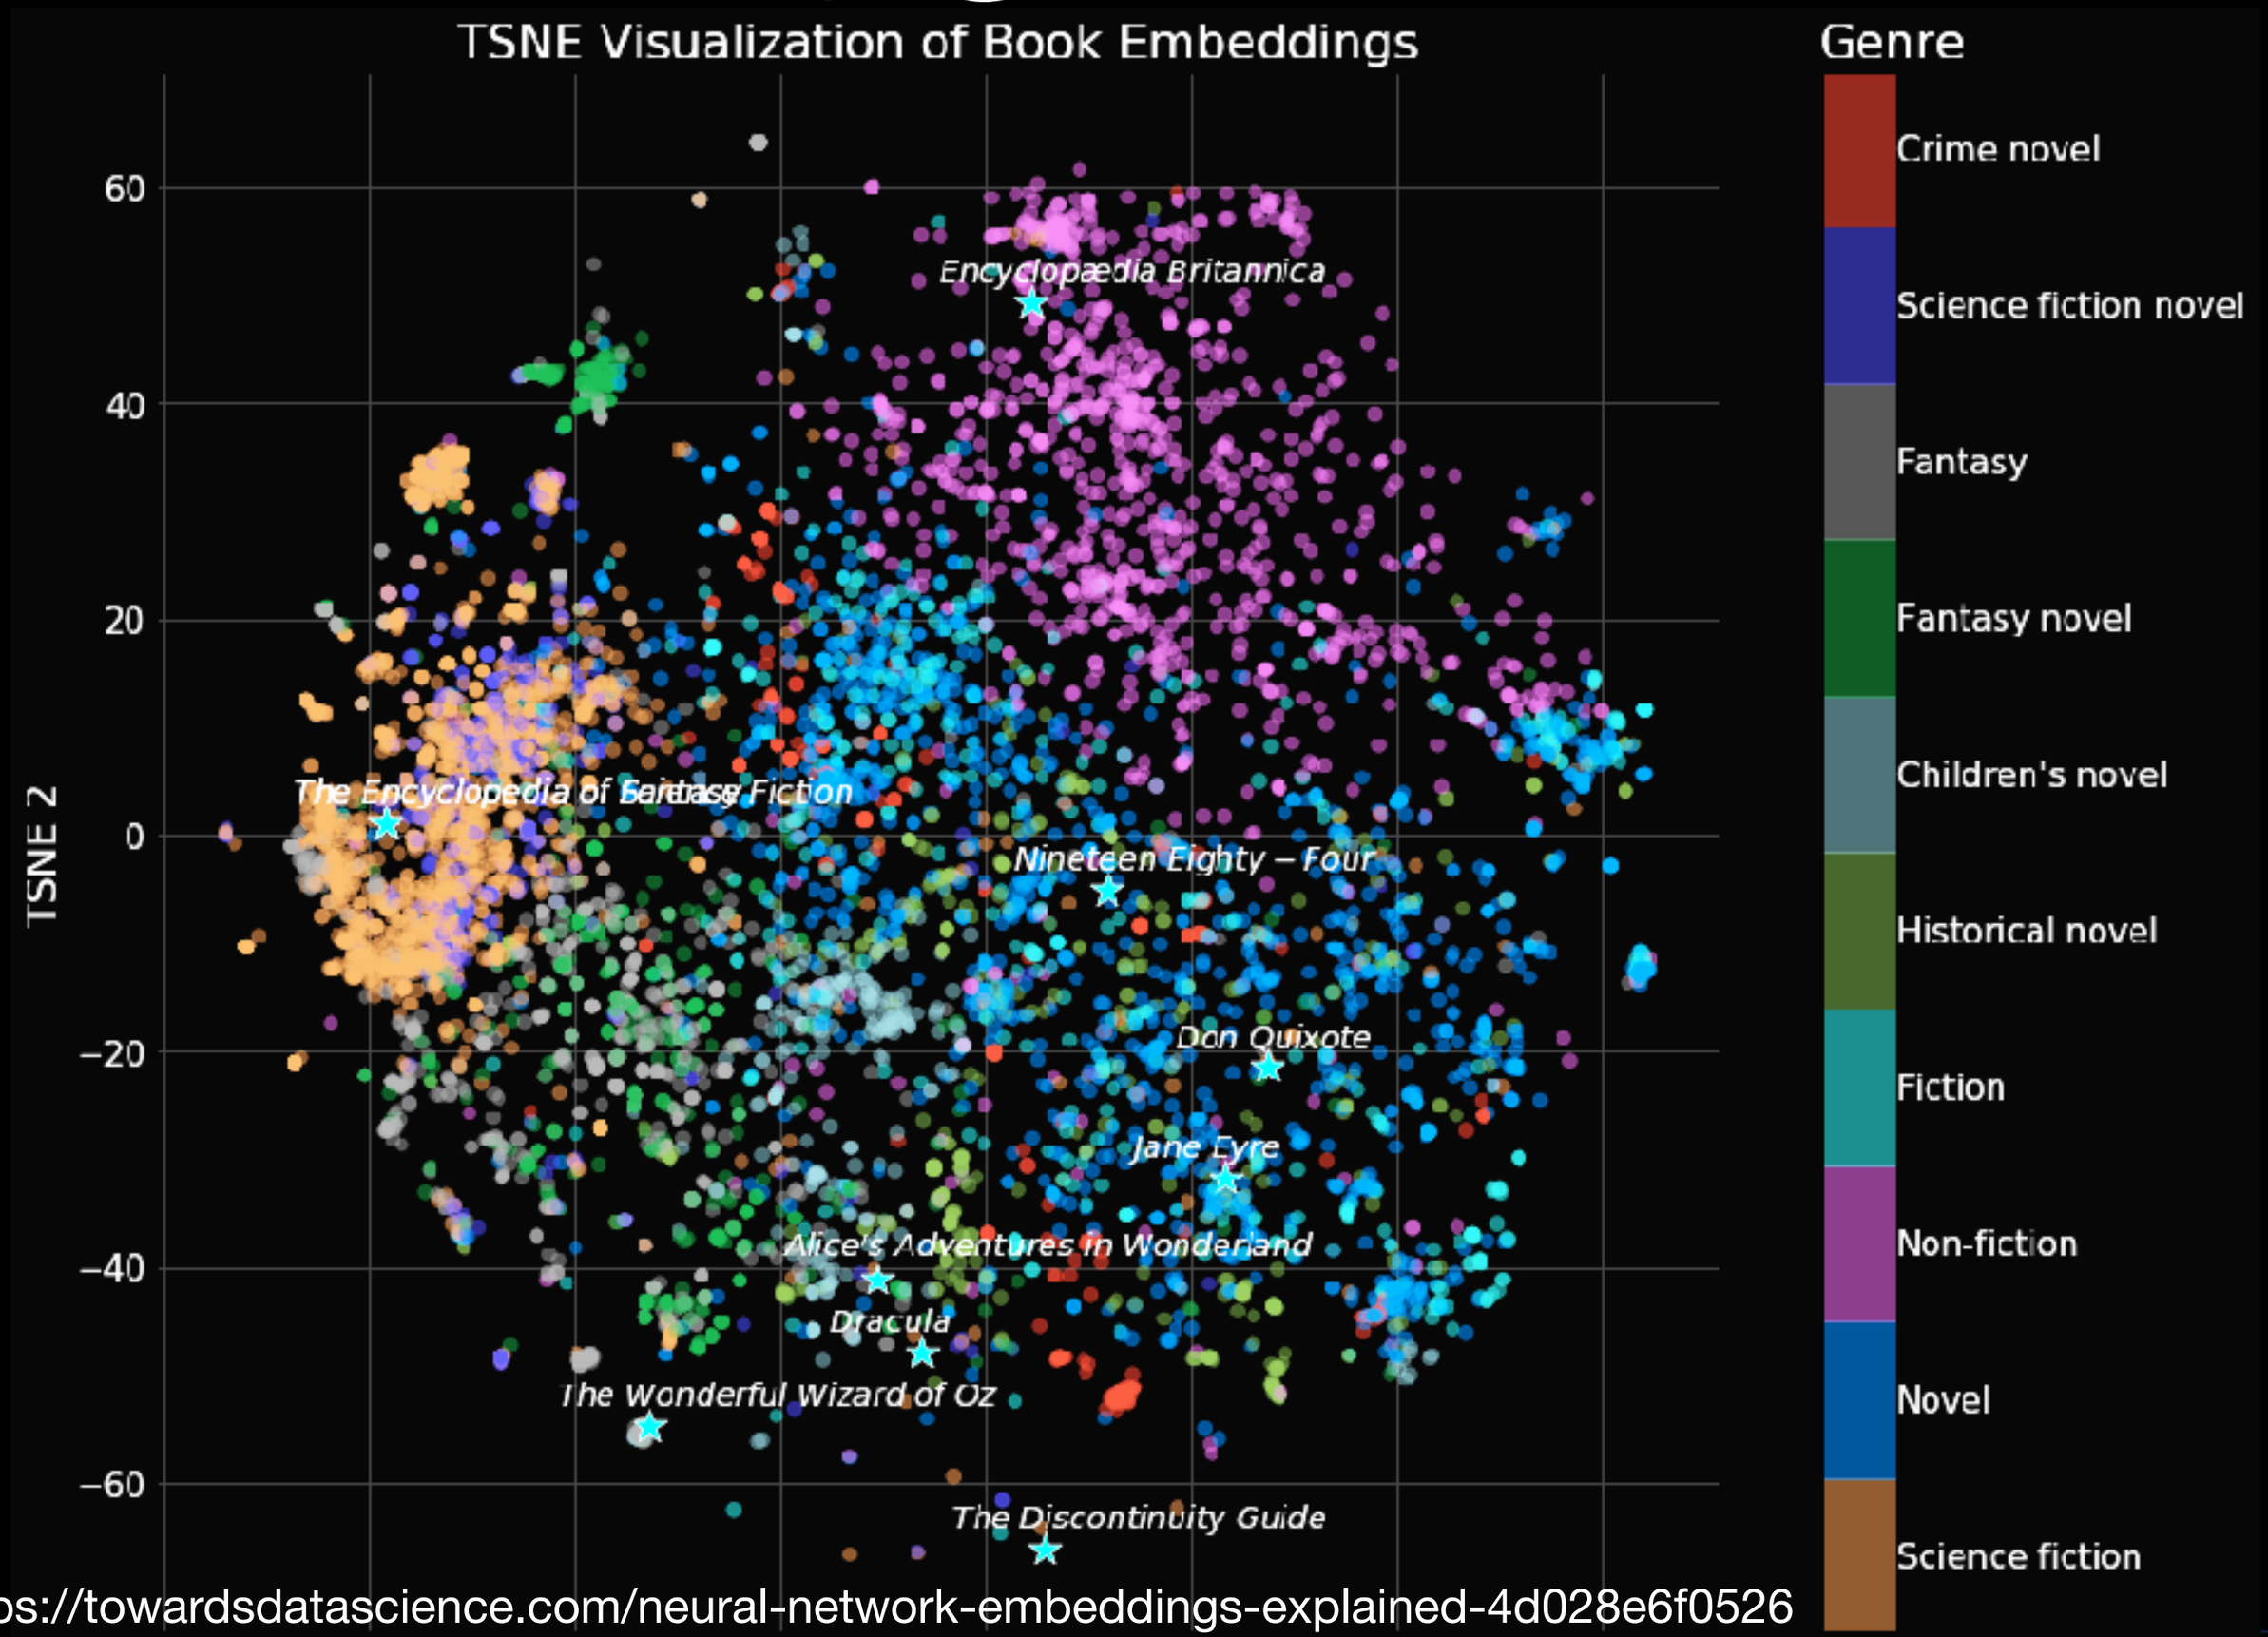

https://towardsdatascience.com/neural-network-embeddings-explained-4d028e6f0526

Clustering

Bocconi

k-Means Clustering

Bocconi

k-Means

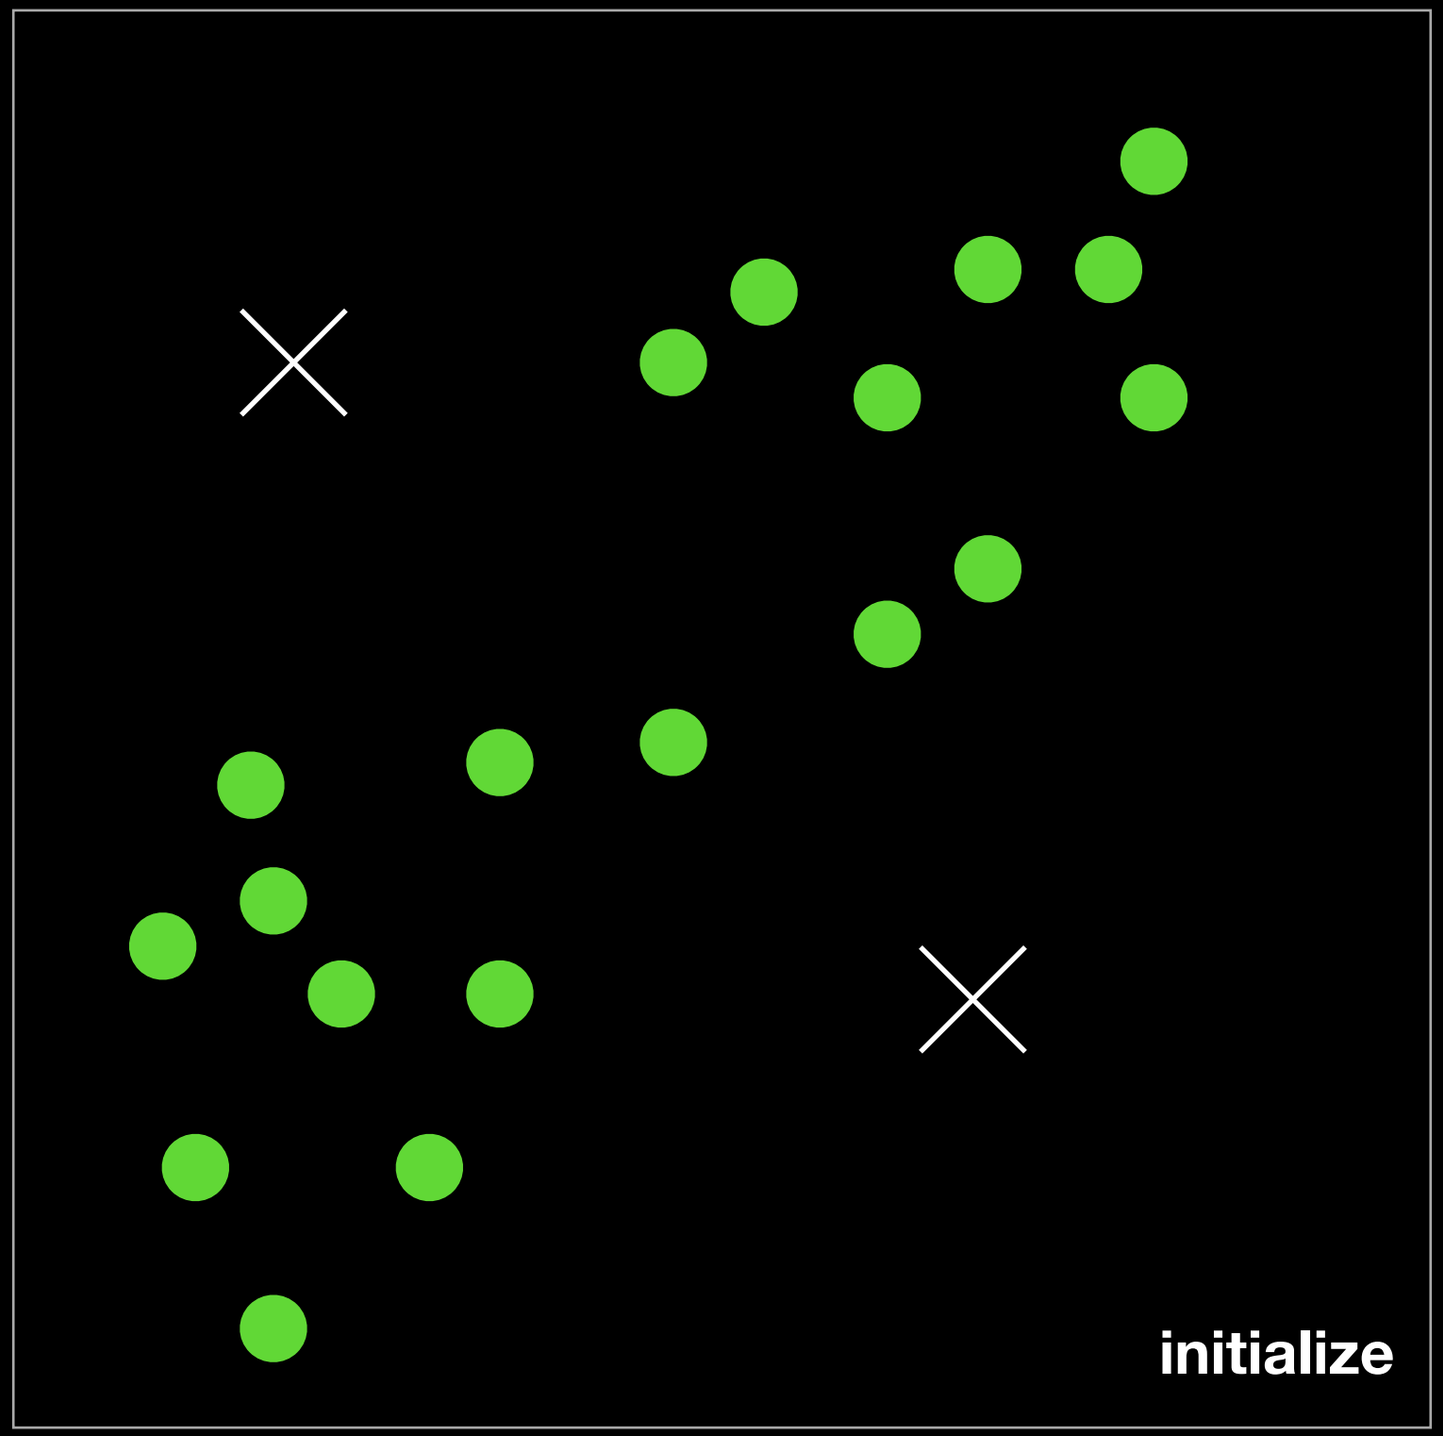

k-Means

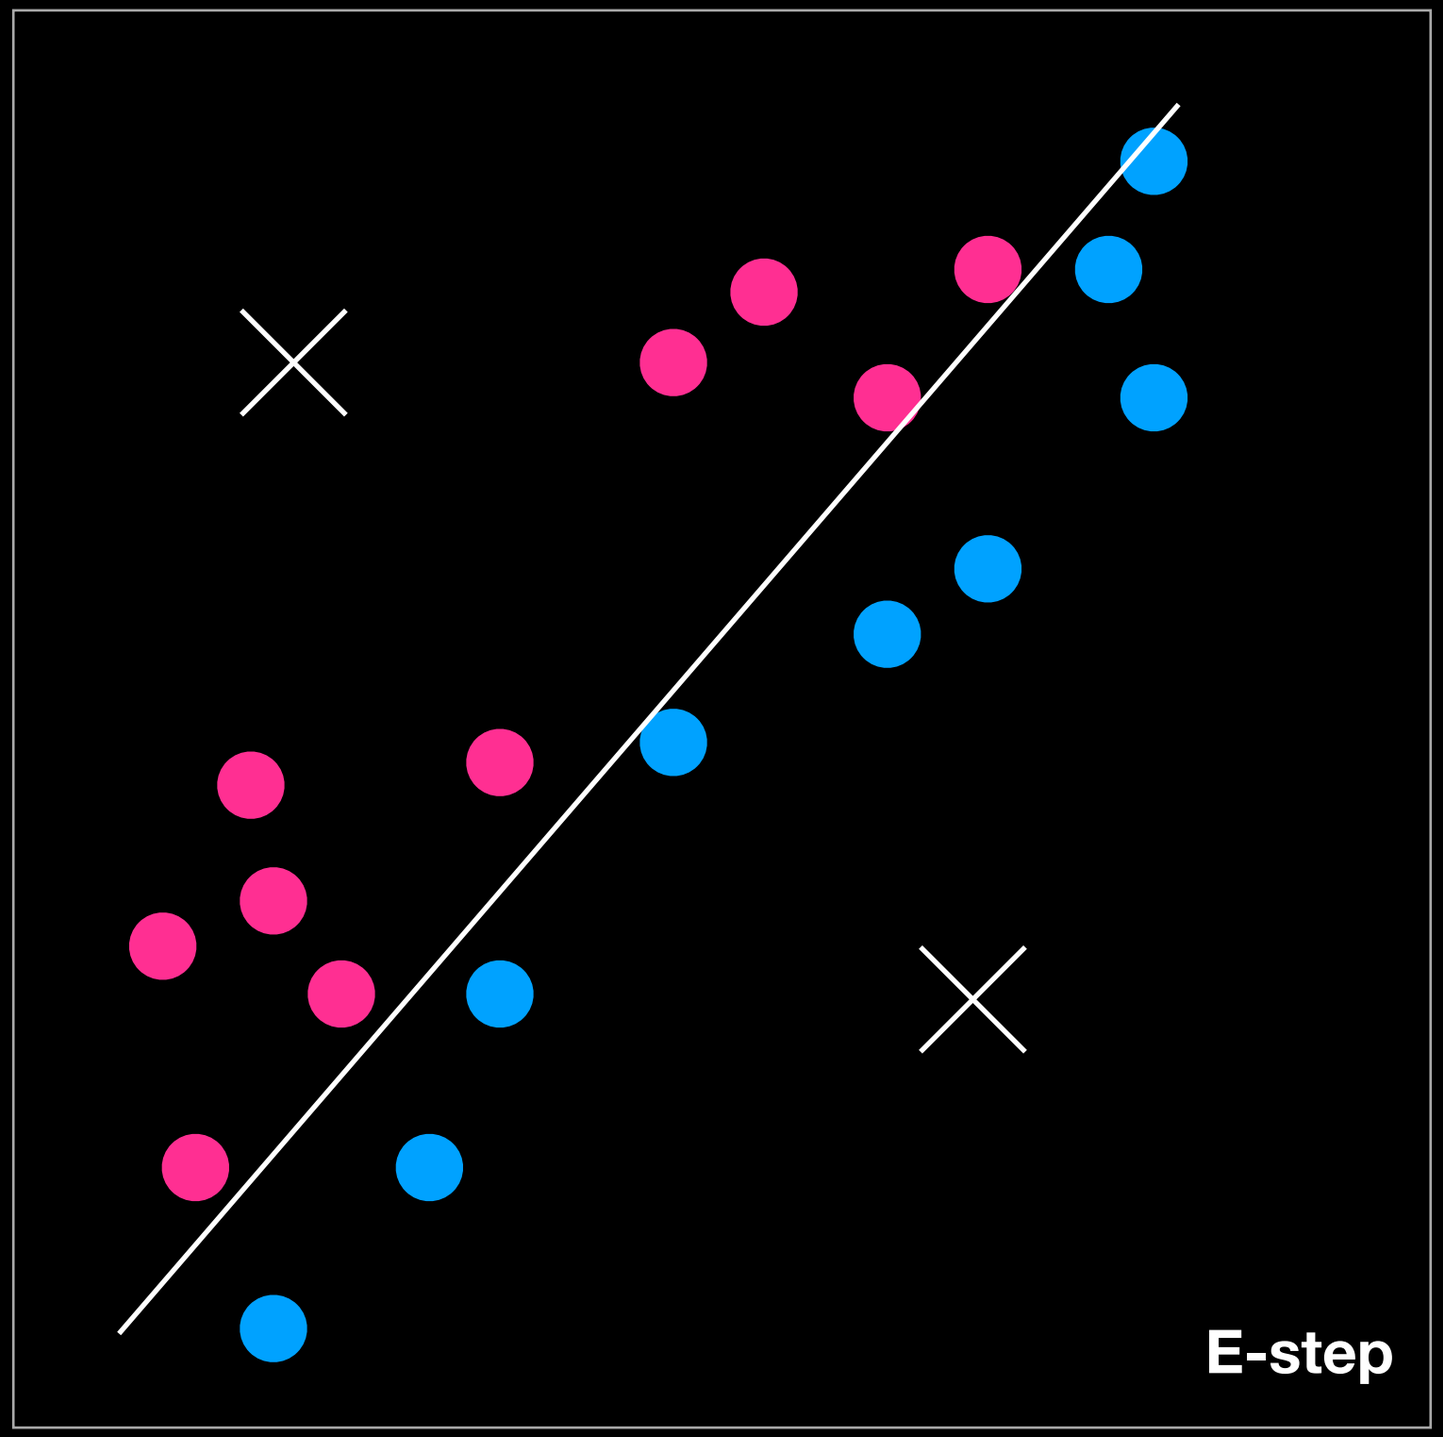

ASSIGN POINTS TO CENTROIDS

k-Means

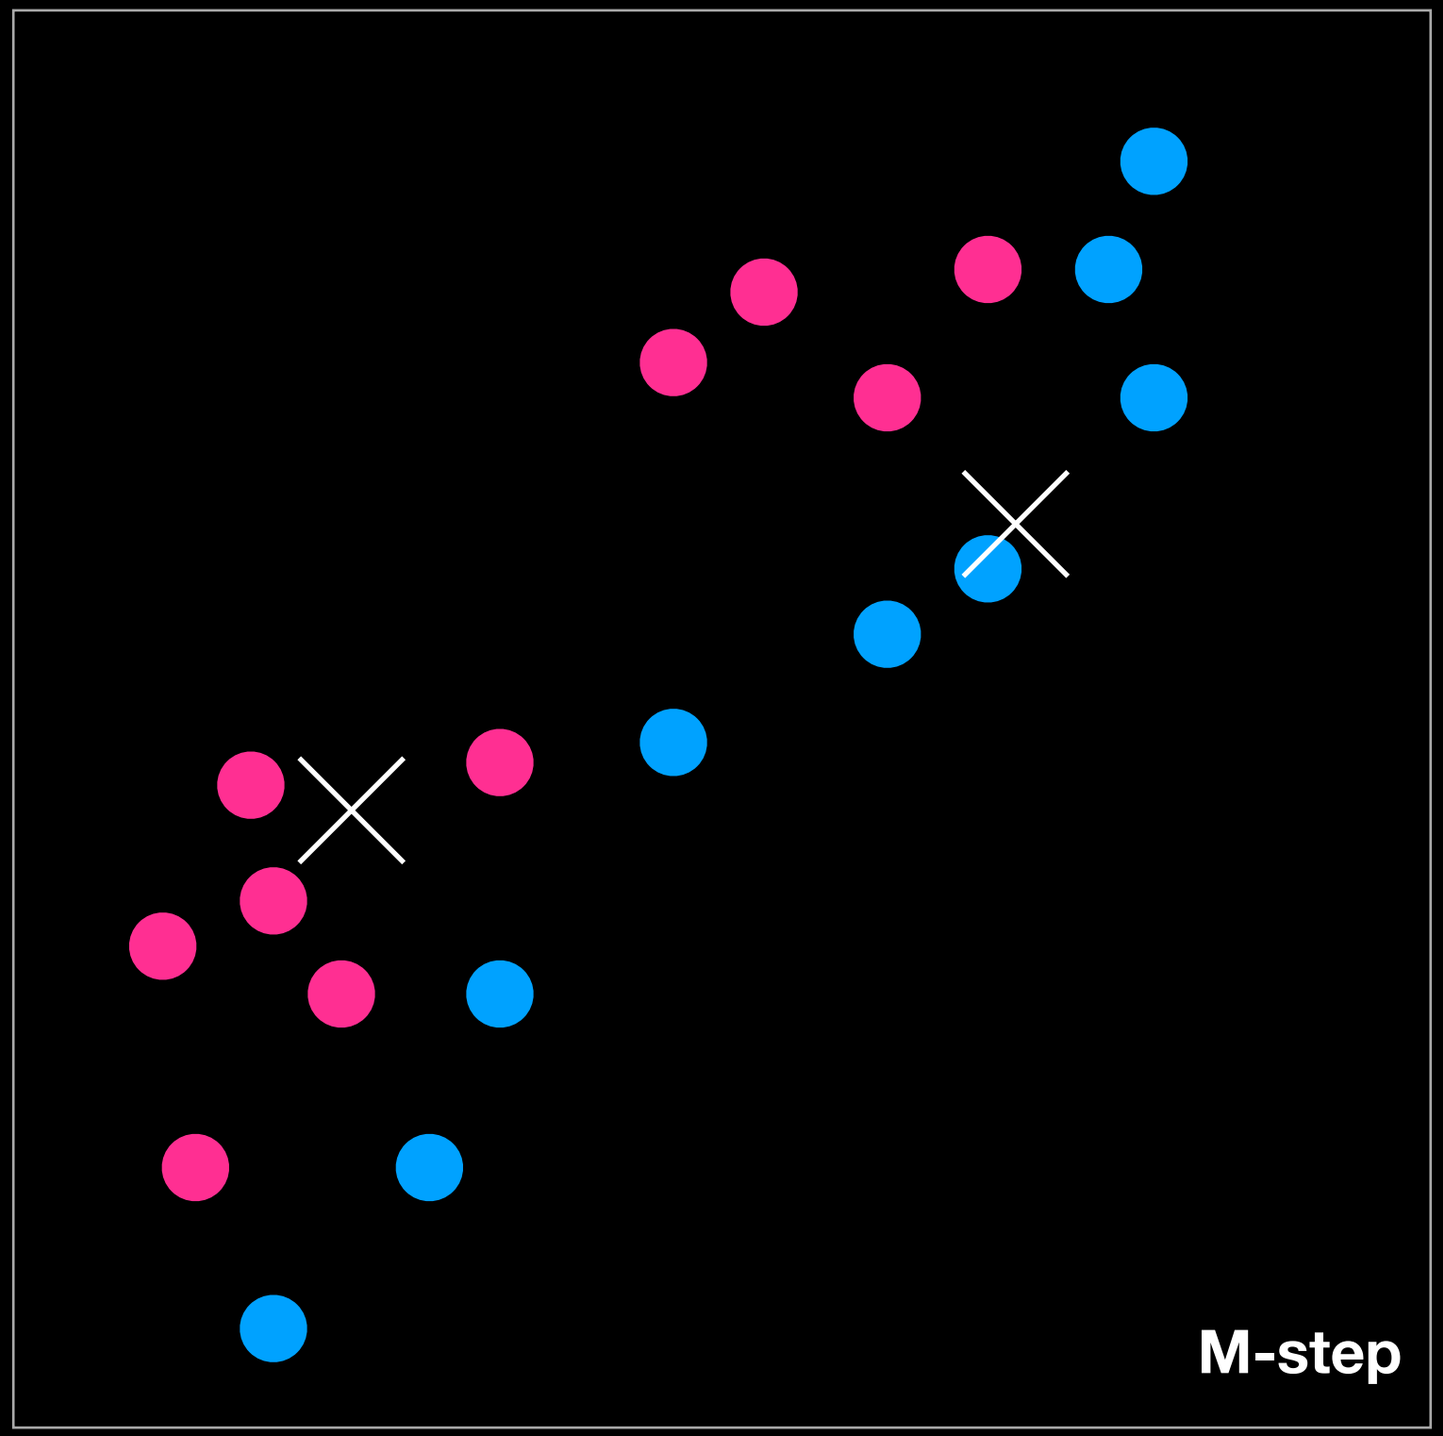

RECOMPUTE CENTROIDS

k-Means

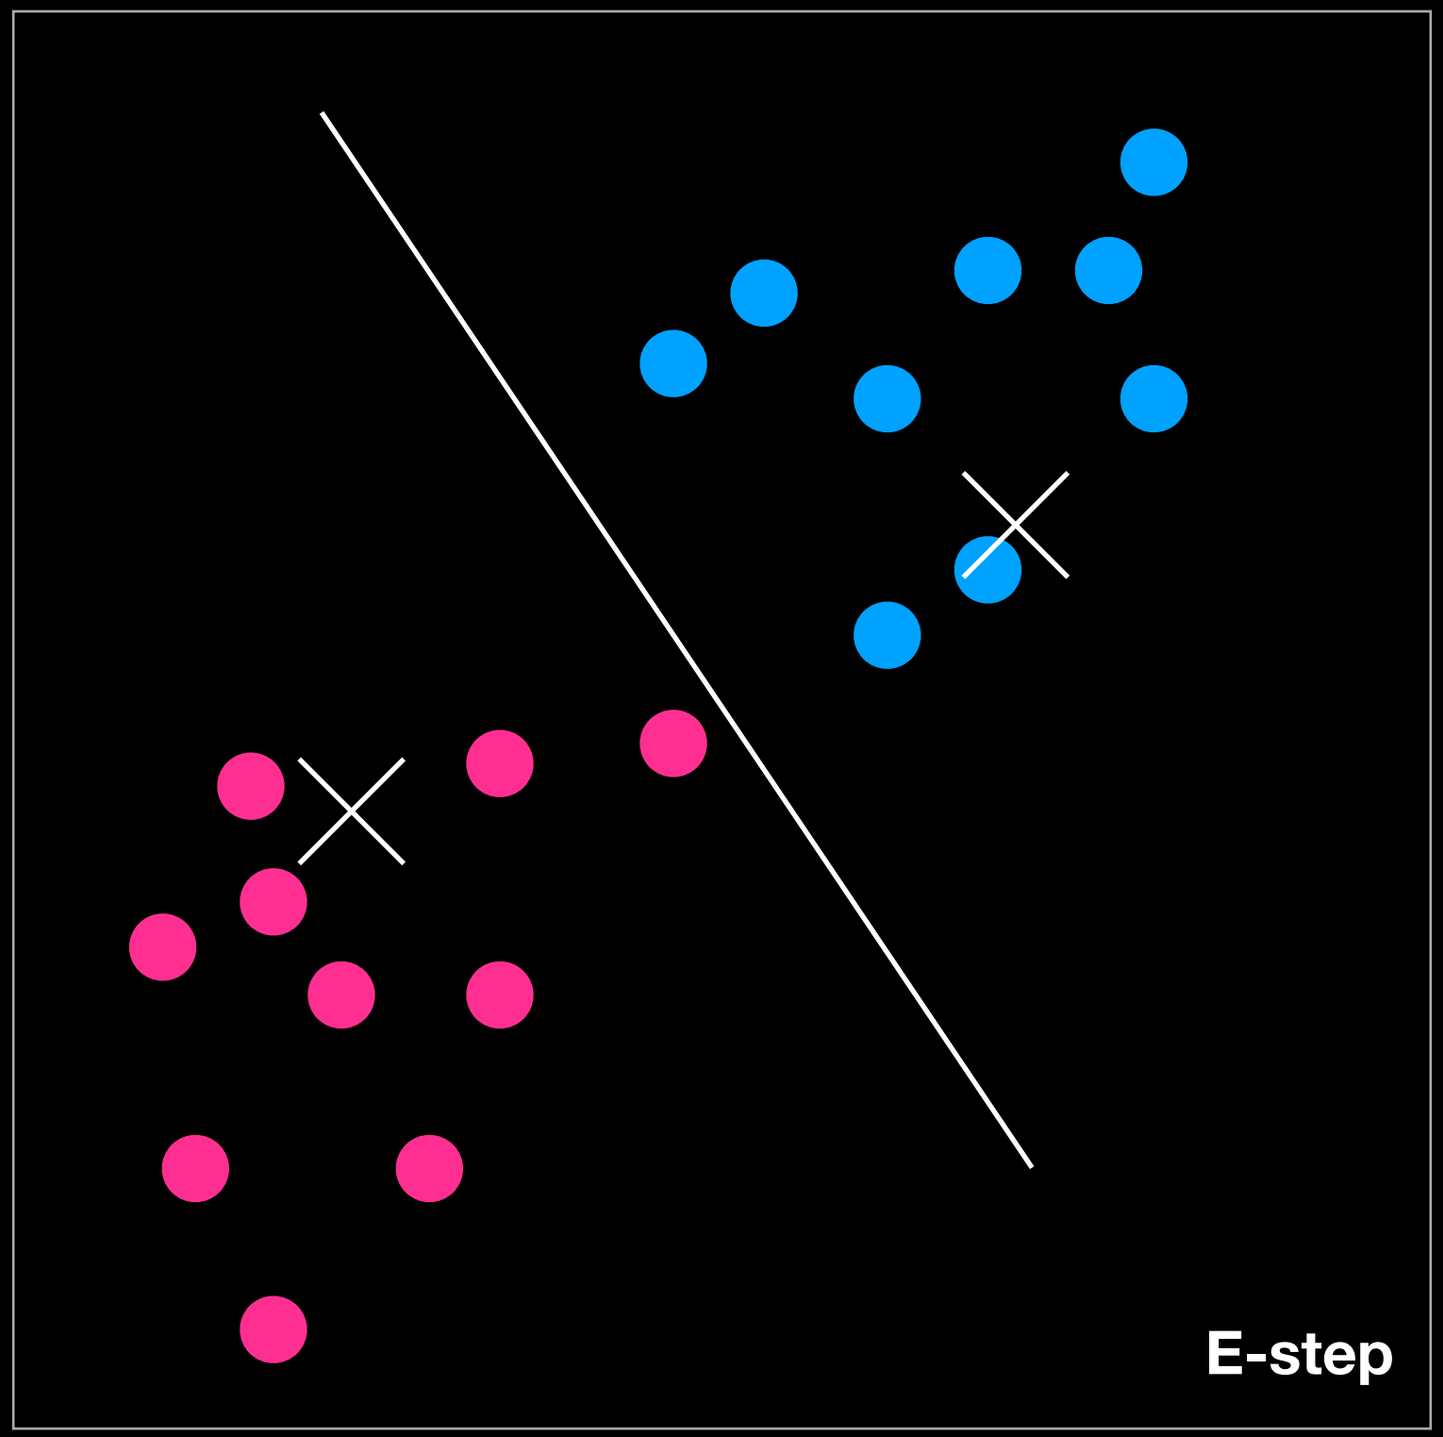

k-Means

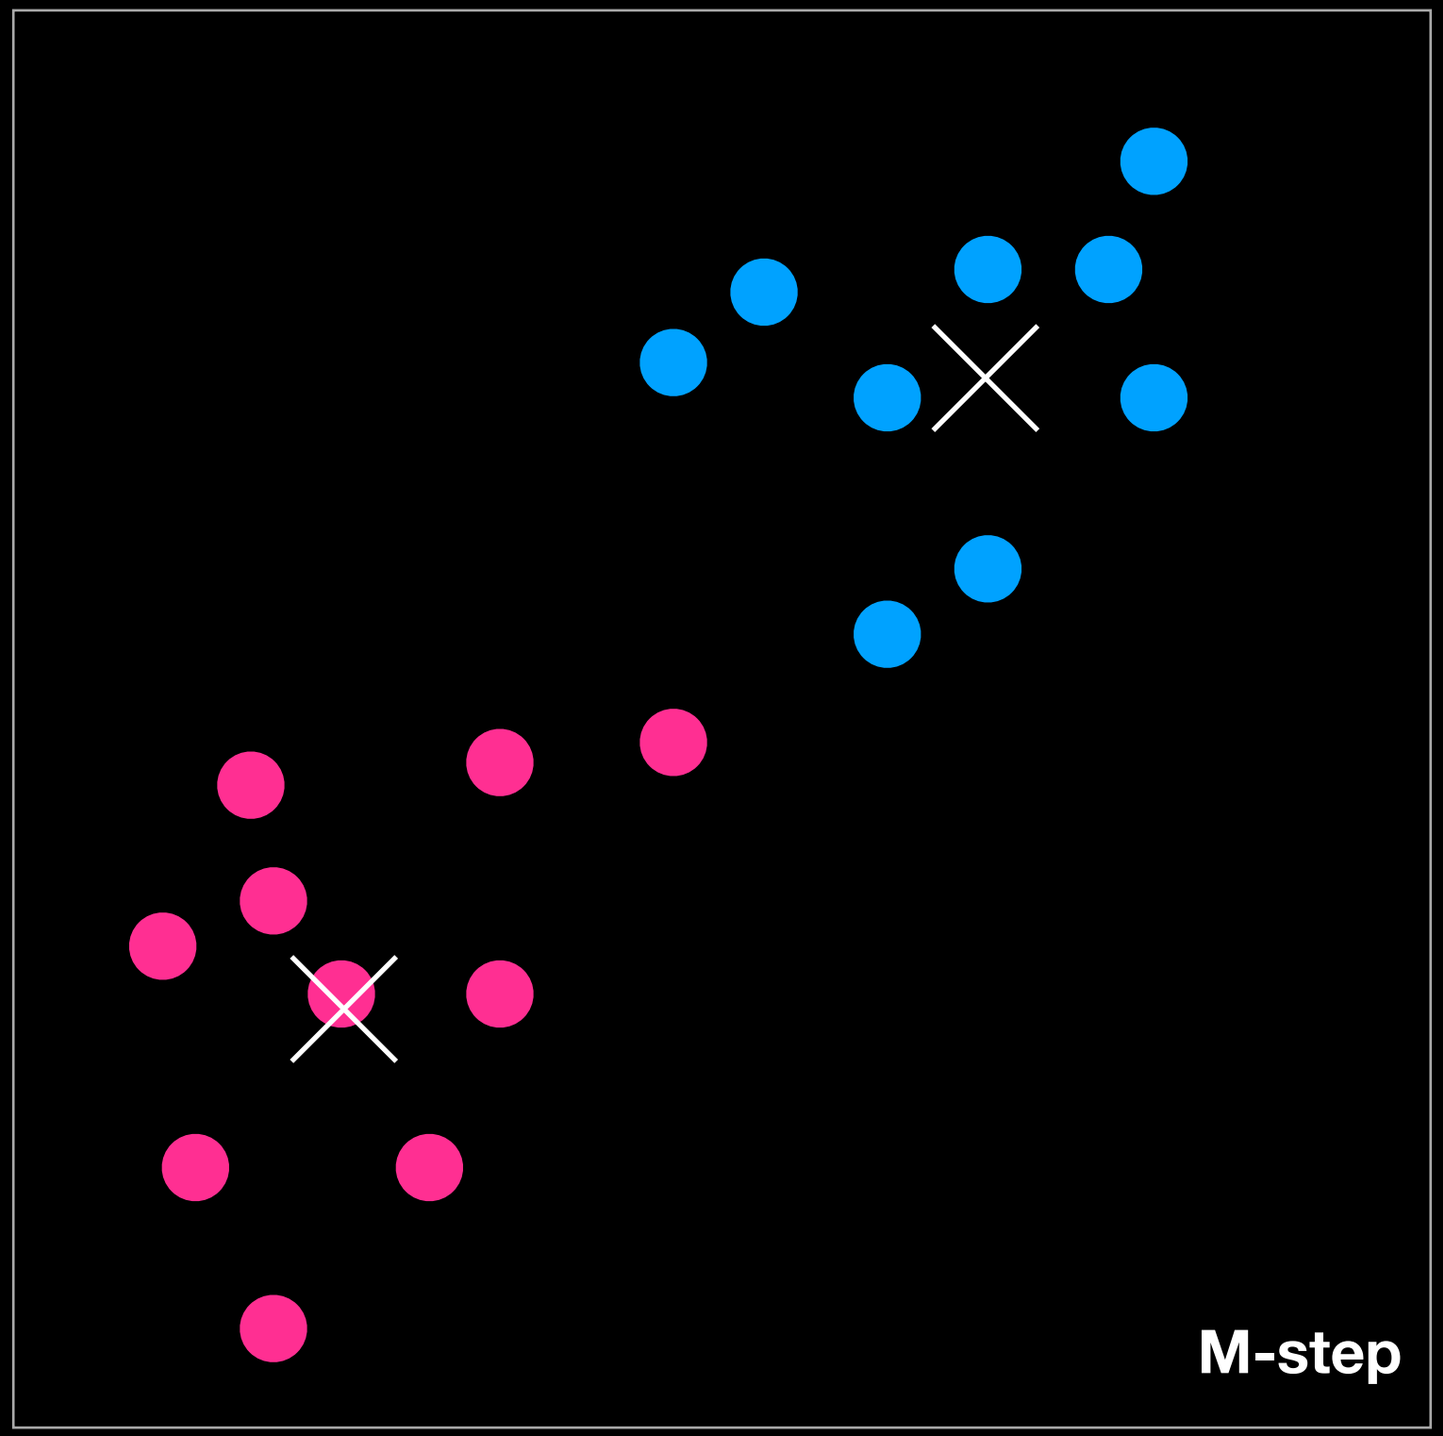

Agglomerative Clustering

Bocconi

Building Up

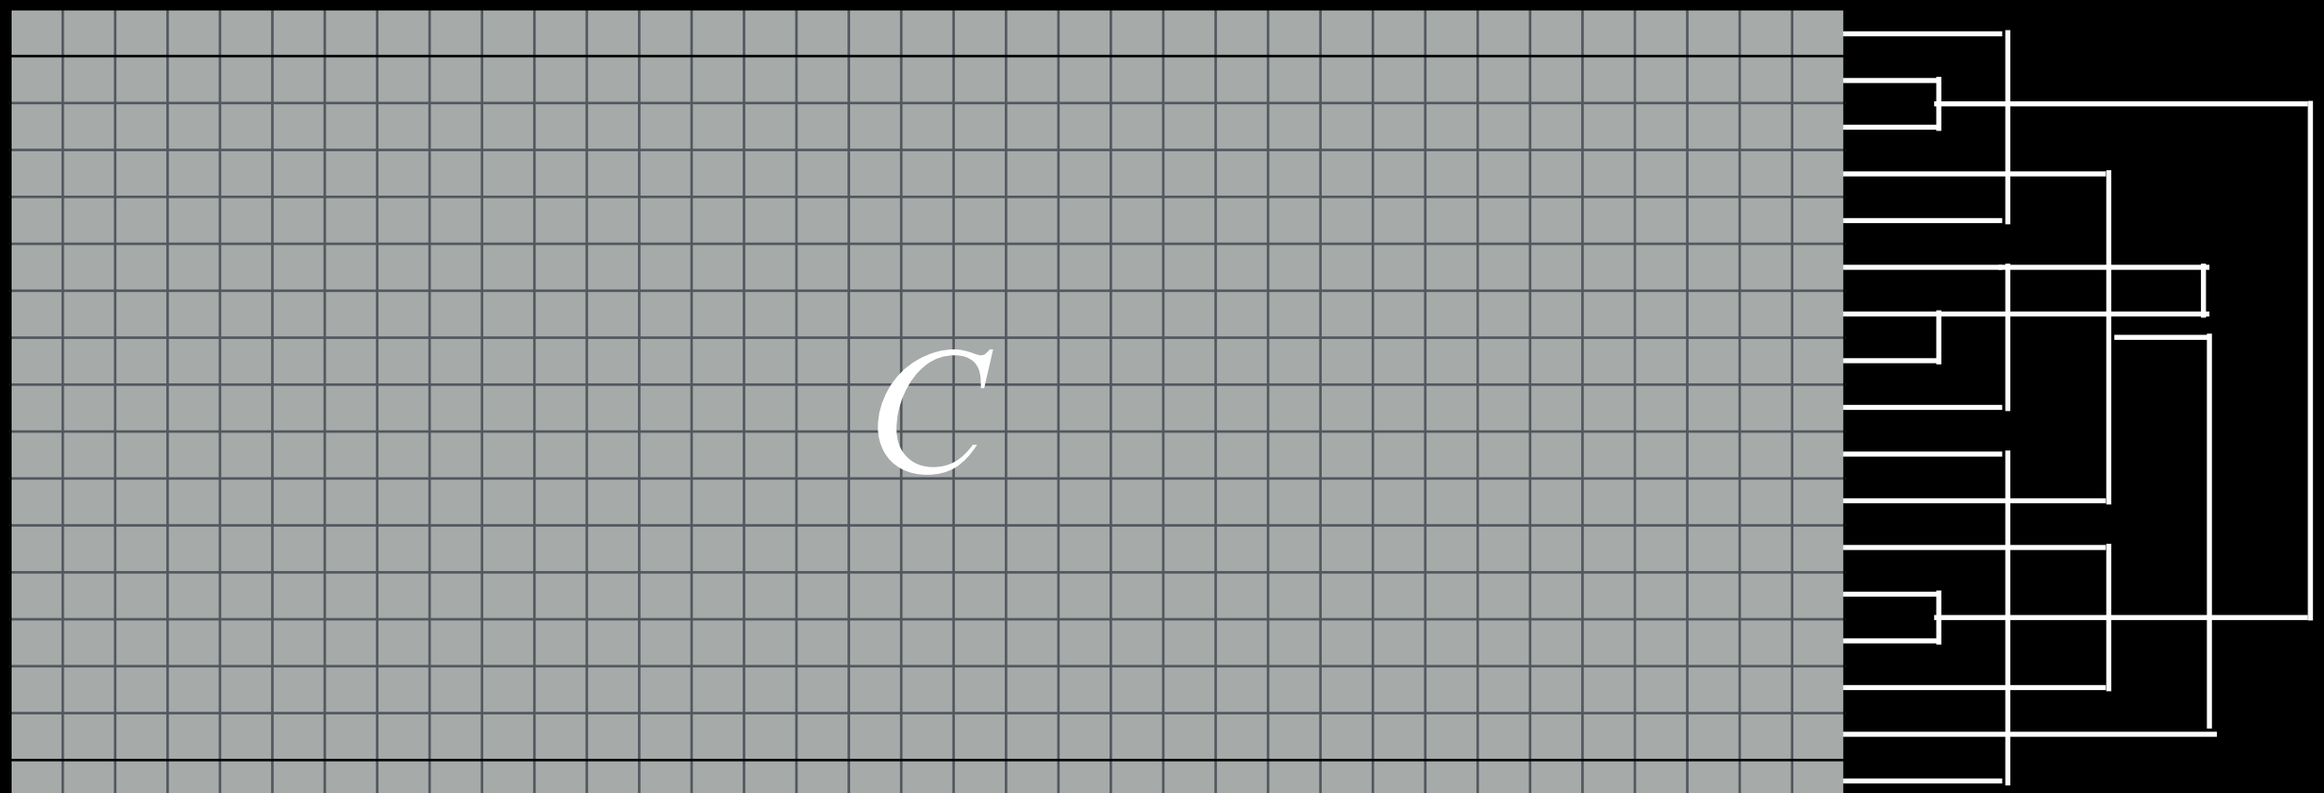

ADJACENCY MATRIX

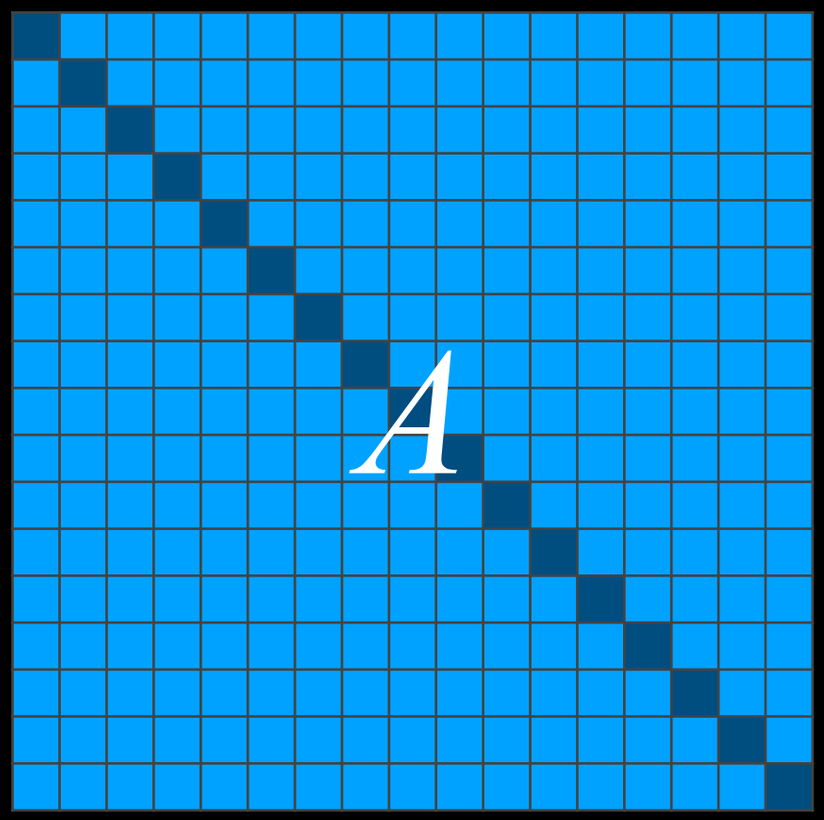

Bocconi

Dialect Clusters

3

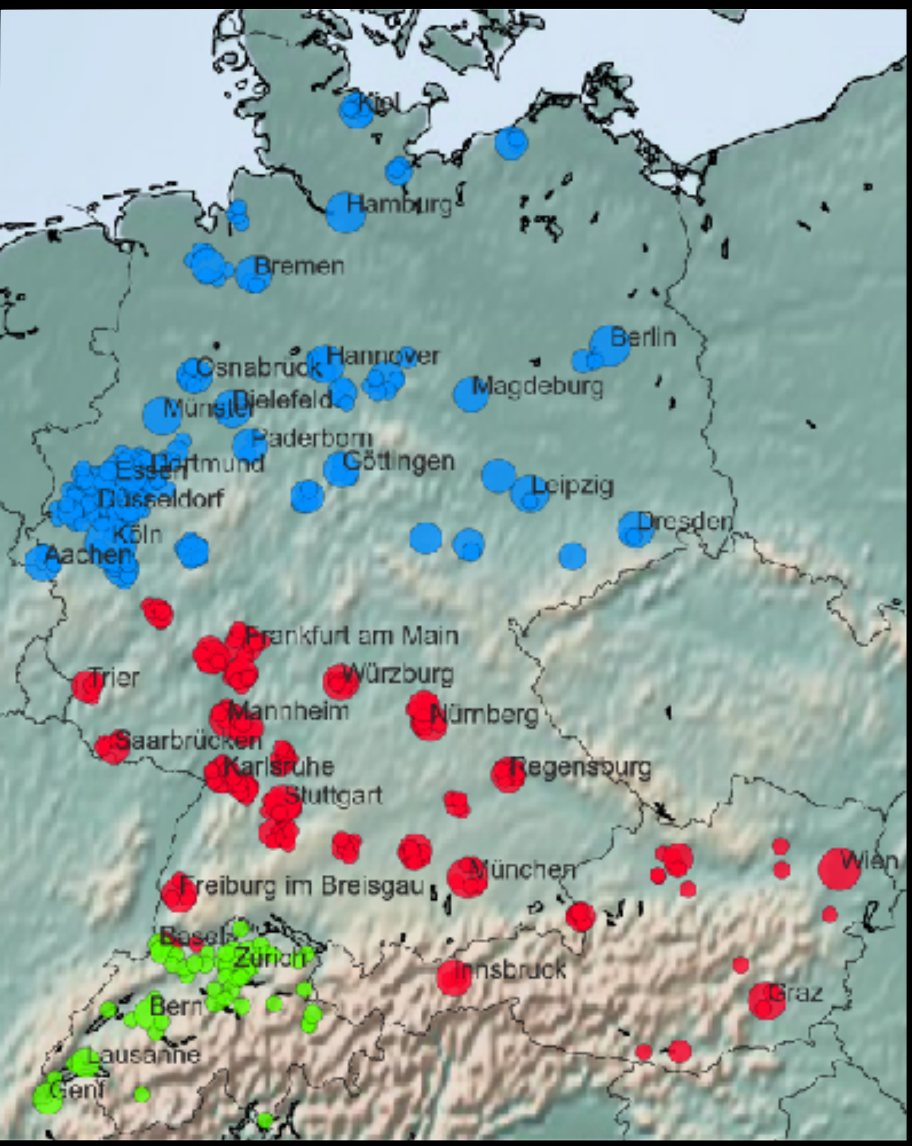

5

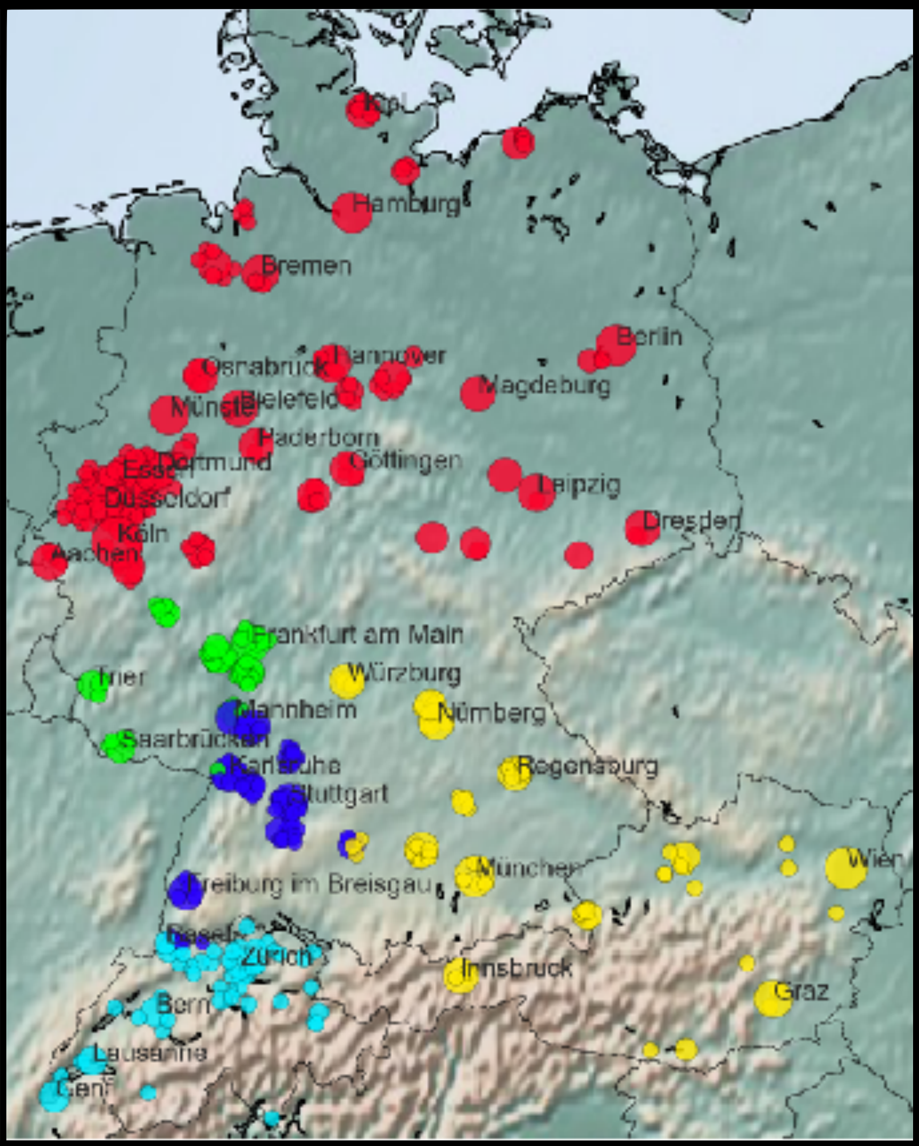

10

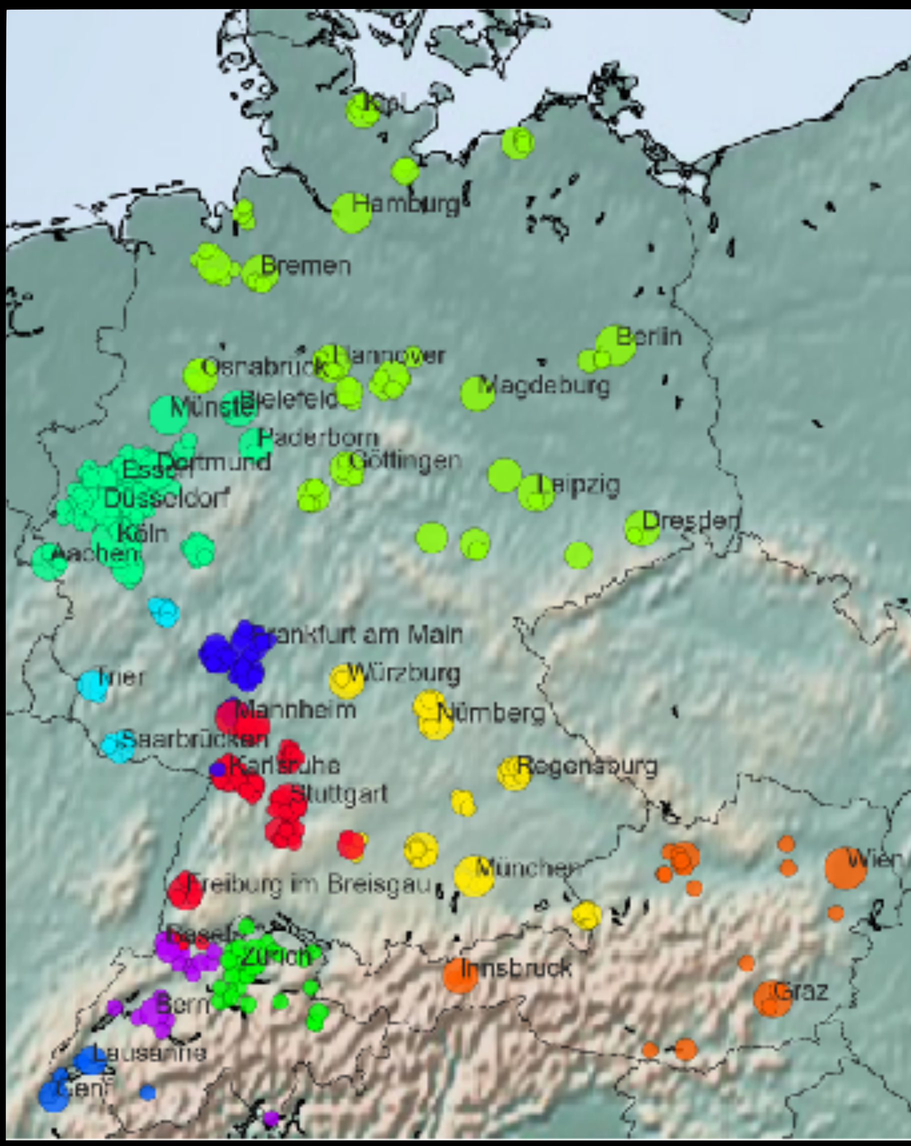

Bocconi

Evaluating Clusters

Bocconi

Making Sense of Clusters

NOUNS

NEUTRAL

VERBS

perspicacious

lacerate

president

kissed

fresh

induce

rotten

ROMANCE ROOT

GERMANIC ROOT

platypus

arsenic

mortify

mortgage

recognize

pie

blue

POSITIVE

ADJECTIVES

NEGATIVE

How Many Clusters?

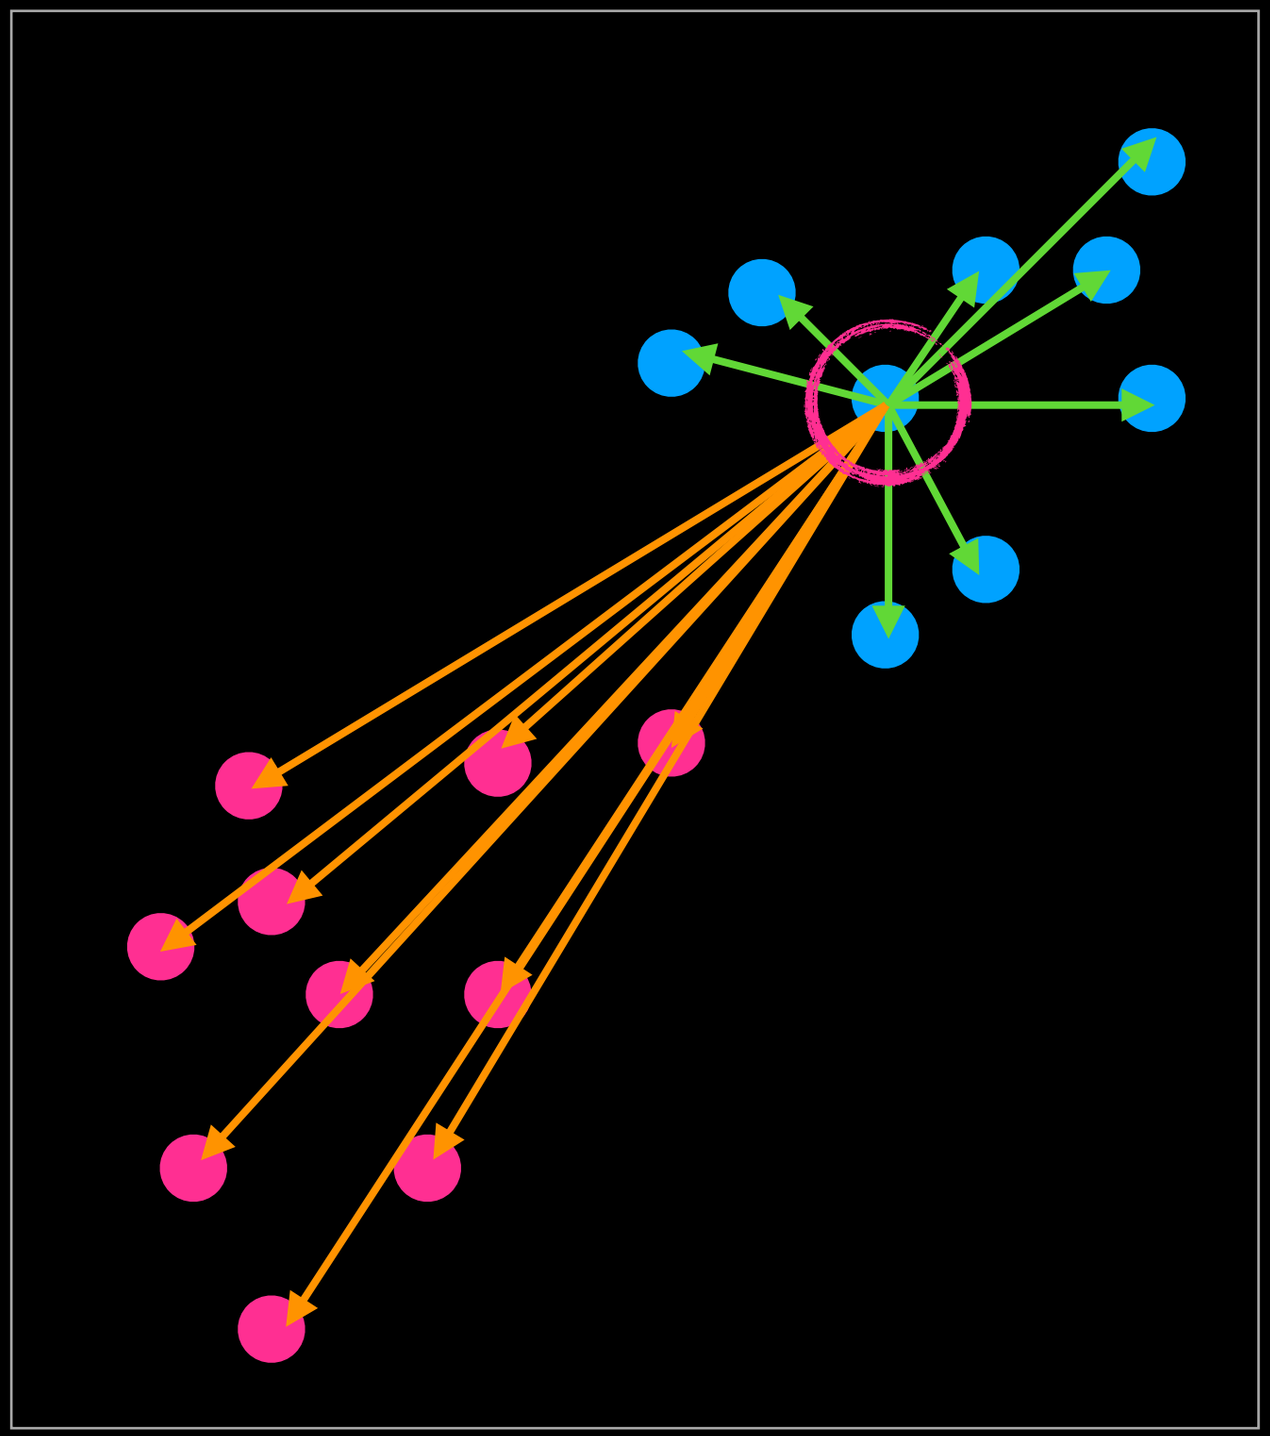

Silhouette Score

a = mean intra-cluster distance

$$S = \frac{(b - a)}{max(a, b)}$$

b = mean dist. nearest cluster

Silhouette Scores

DEPENDS ON PATIENCE|COMPUTE POWER

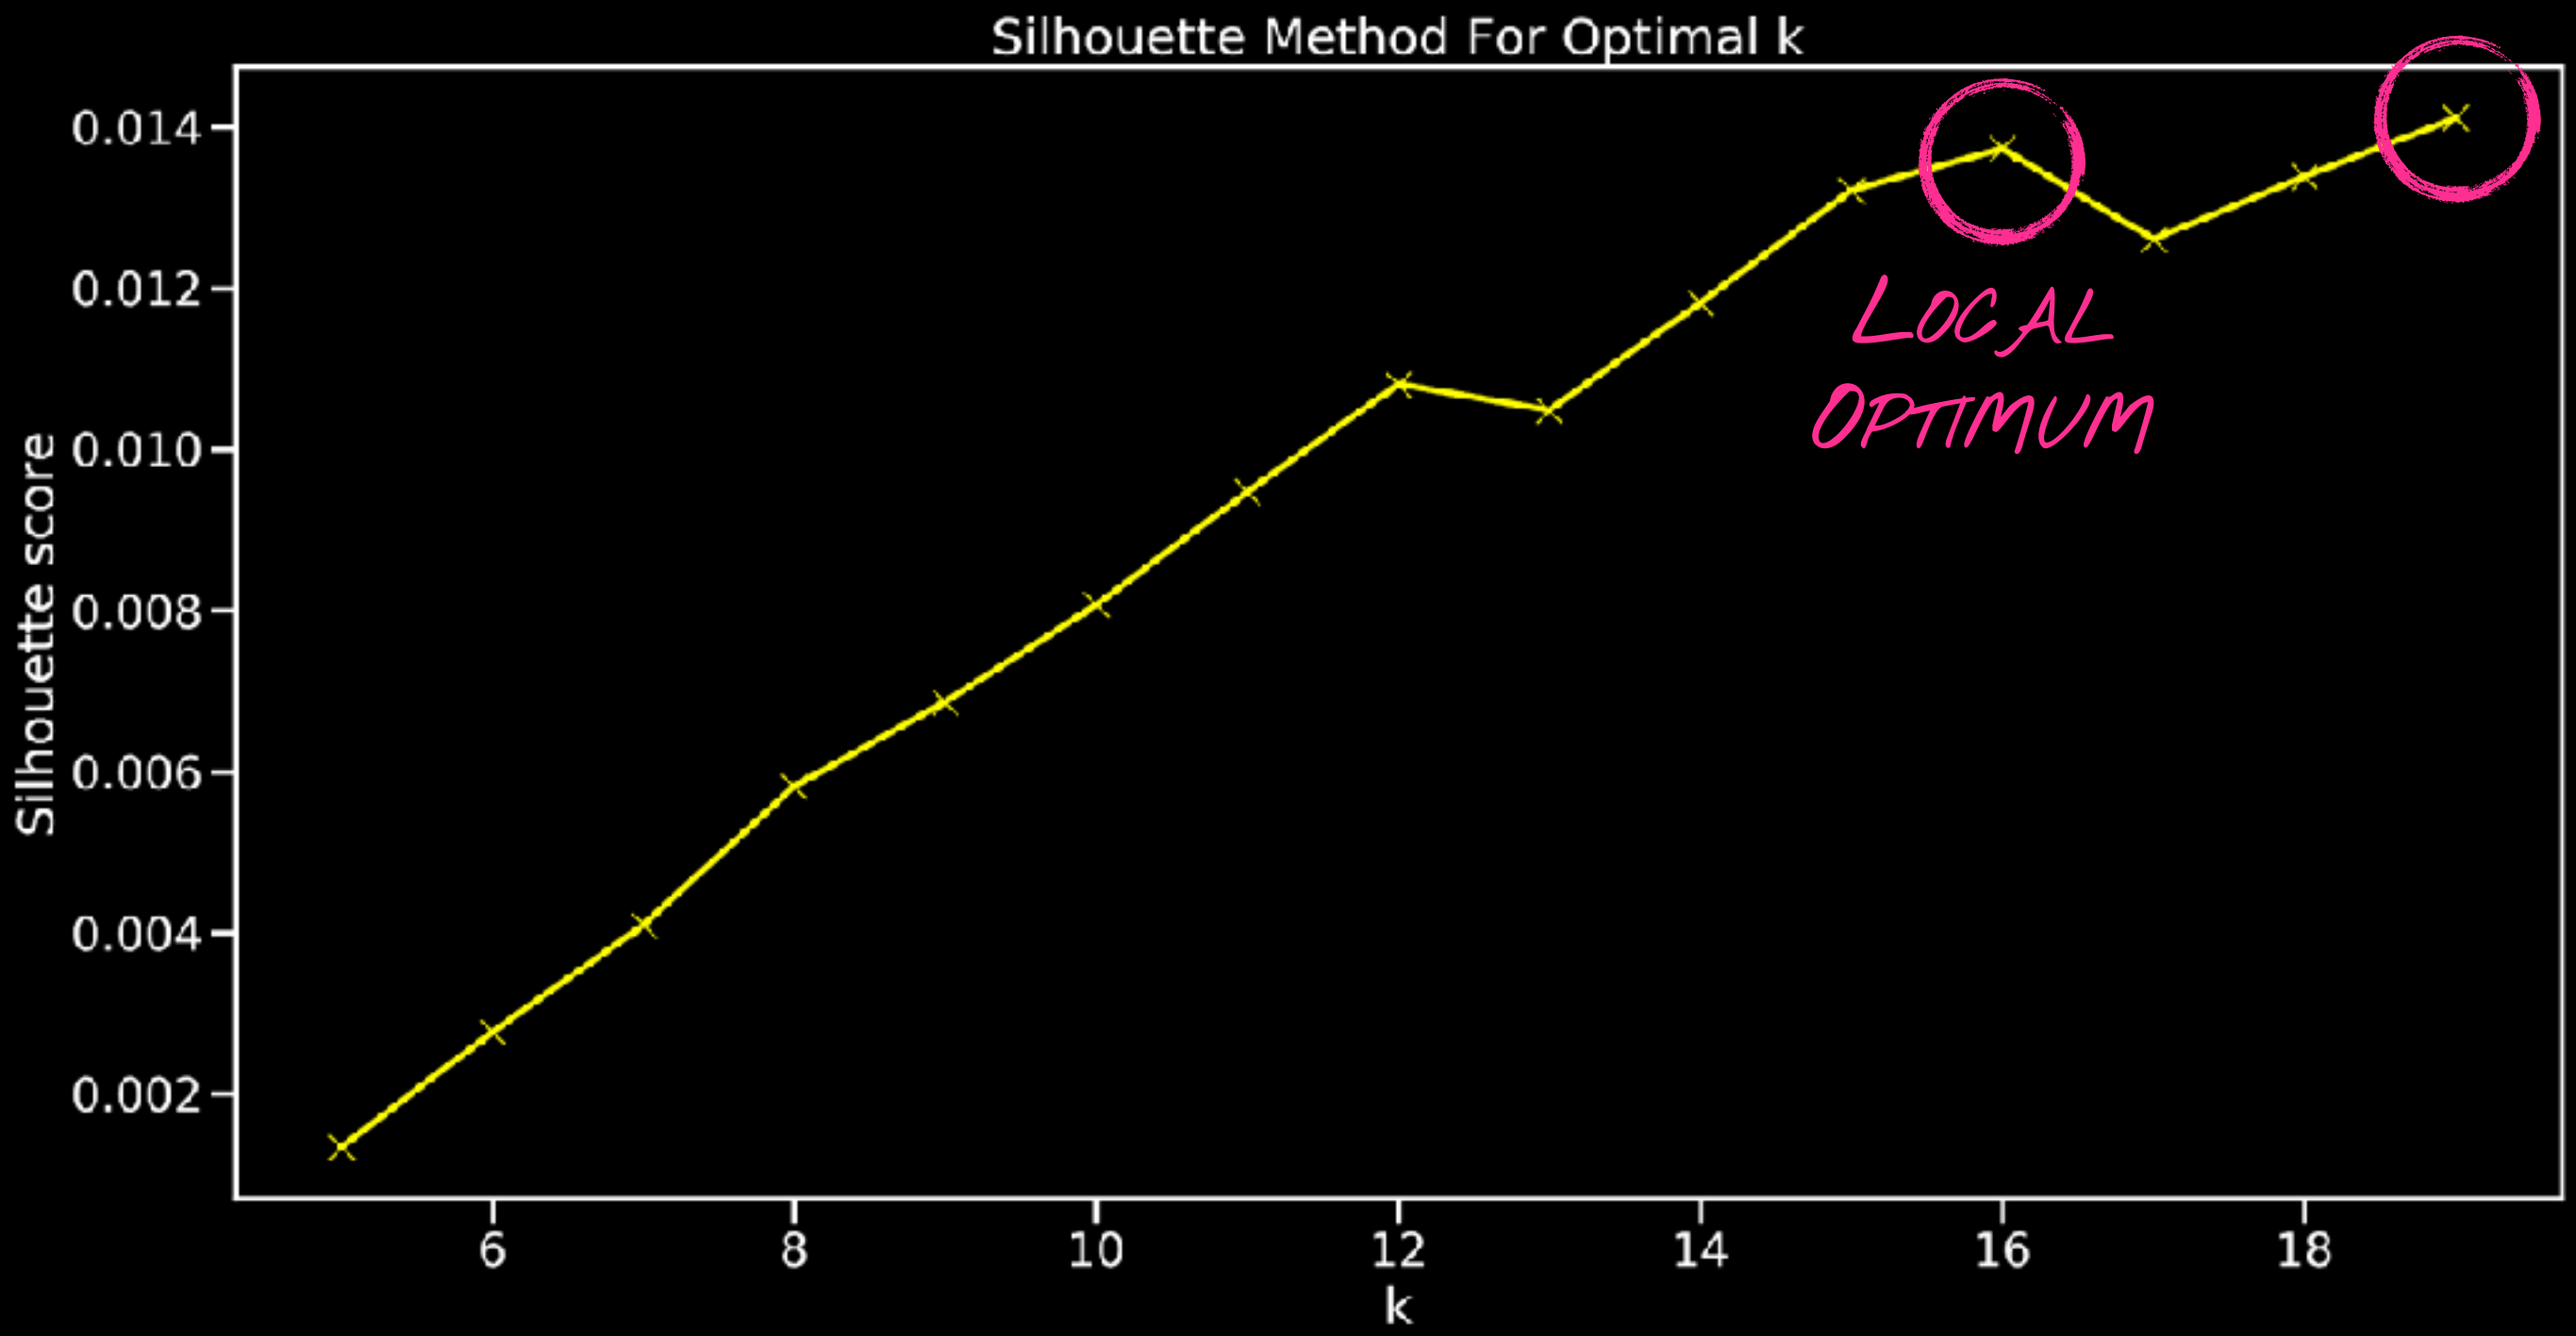

Bocconi

Supervised Evaluation Metrics

	Homogeneity *cluster* has only 1 gold label	**Completeness** *gold label* has only 1 cluster
Good:	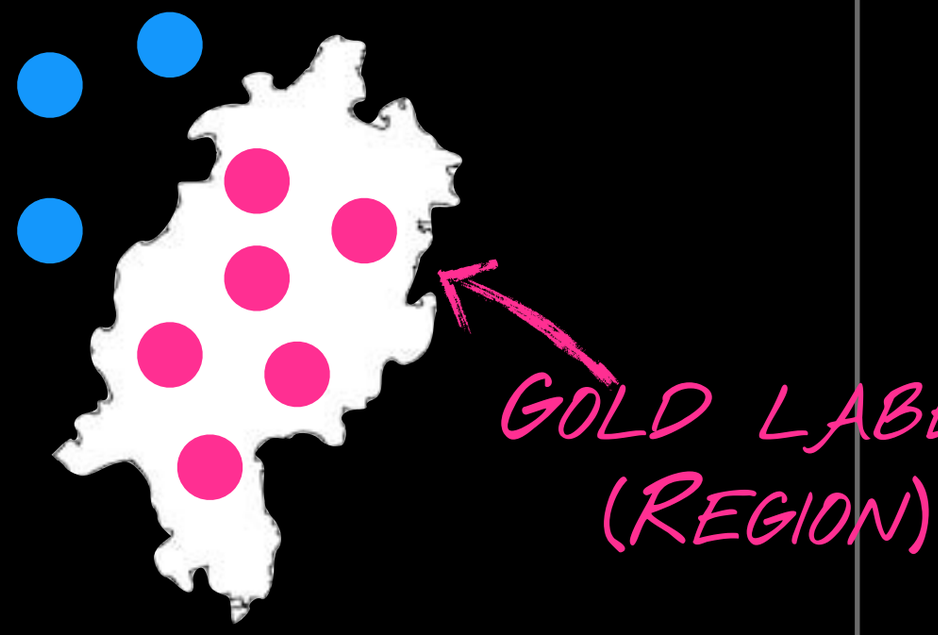	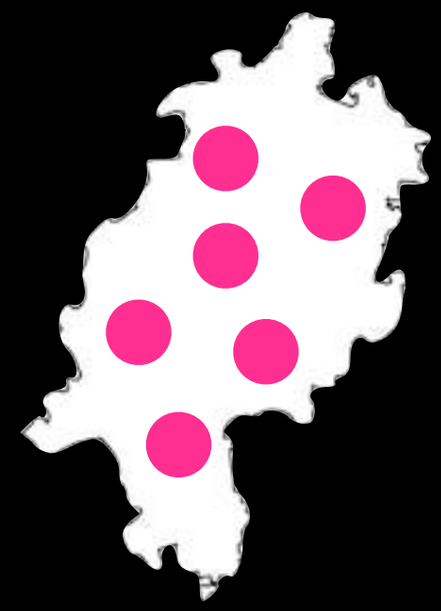
Bad:	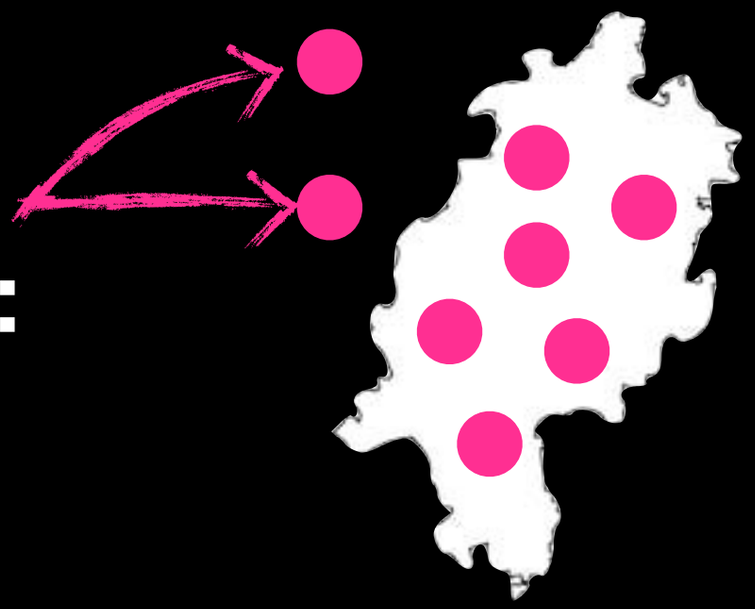	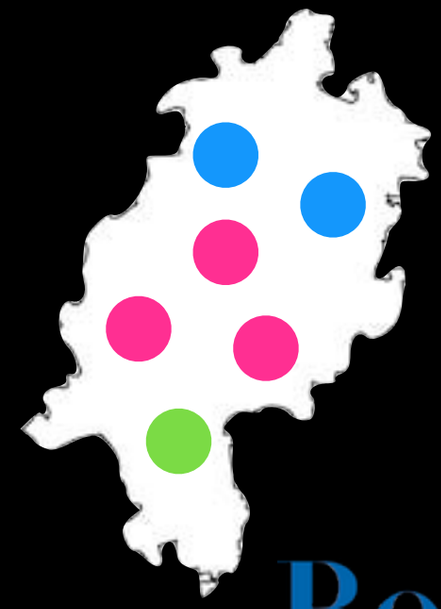

Comparison

	***k*-means**	**Agg**
scalable	yes	no (up to ~20k)
repeatable result	no	yes
include external info	no	yes
Good on dense clusters?	no	yes

Wrapping Up

Bocconi

When to Use What

	Discrete Features	Embeddings
Latent topics	NMF	*Not applicable*
RGB translation	NMF	SVD + scaling
Plotting	SVD	t-SNE
Clustering	*Reduce dimensions*	*Use as-is*

Take-Home Points

- **Matrix factorization** assumes latent concept dimensions
  - Can be used for semantic similarity (**LSA**)
  - Reduced components can be **visualized** in **graphs** or as **RGB** colors
- **Clusters** can group input in new ways
- Trade-off between speed and interpretability